

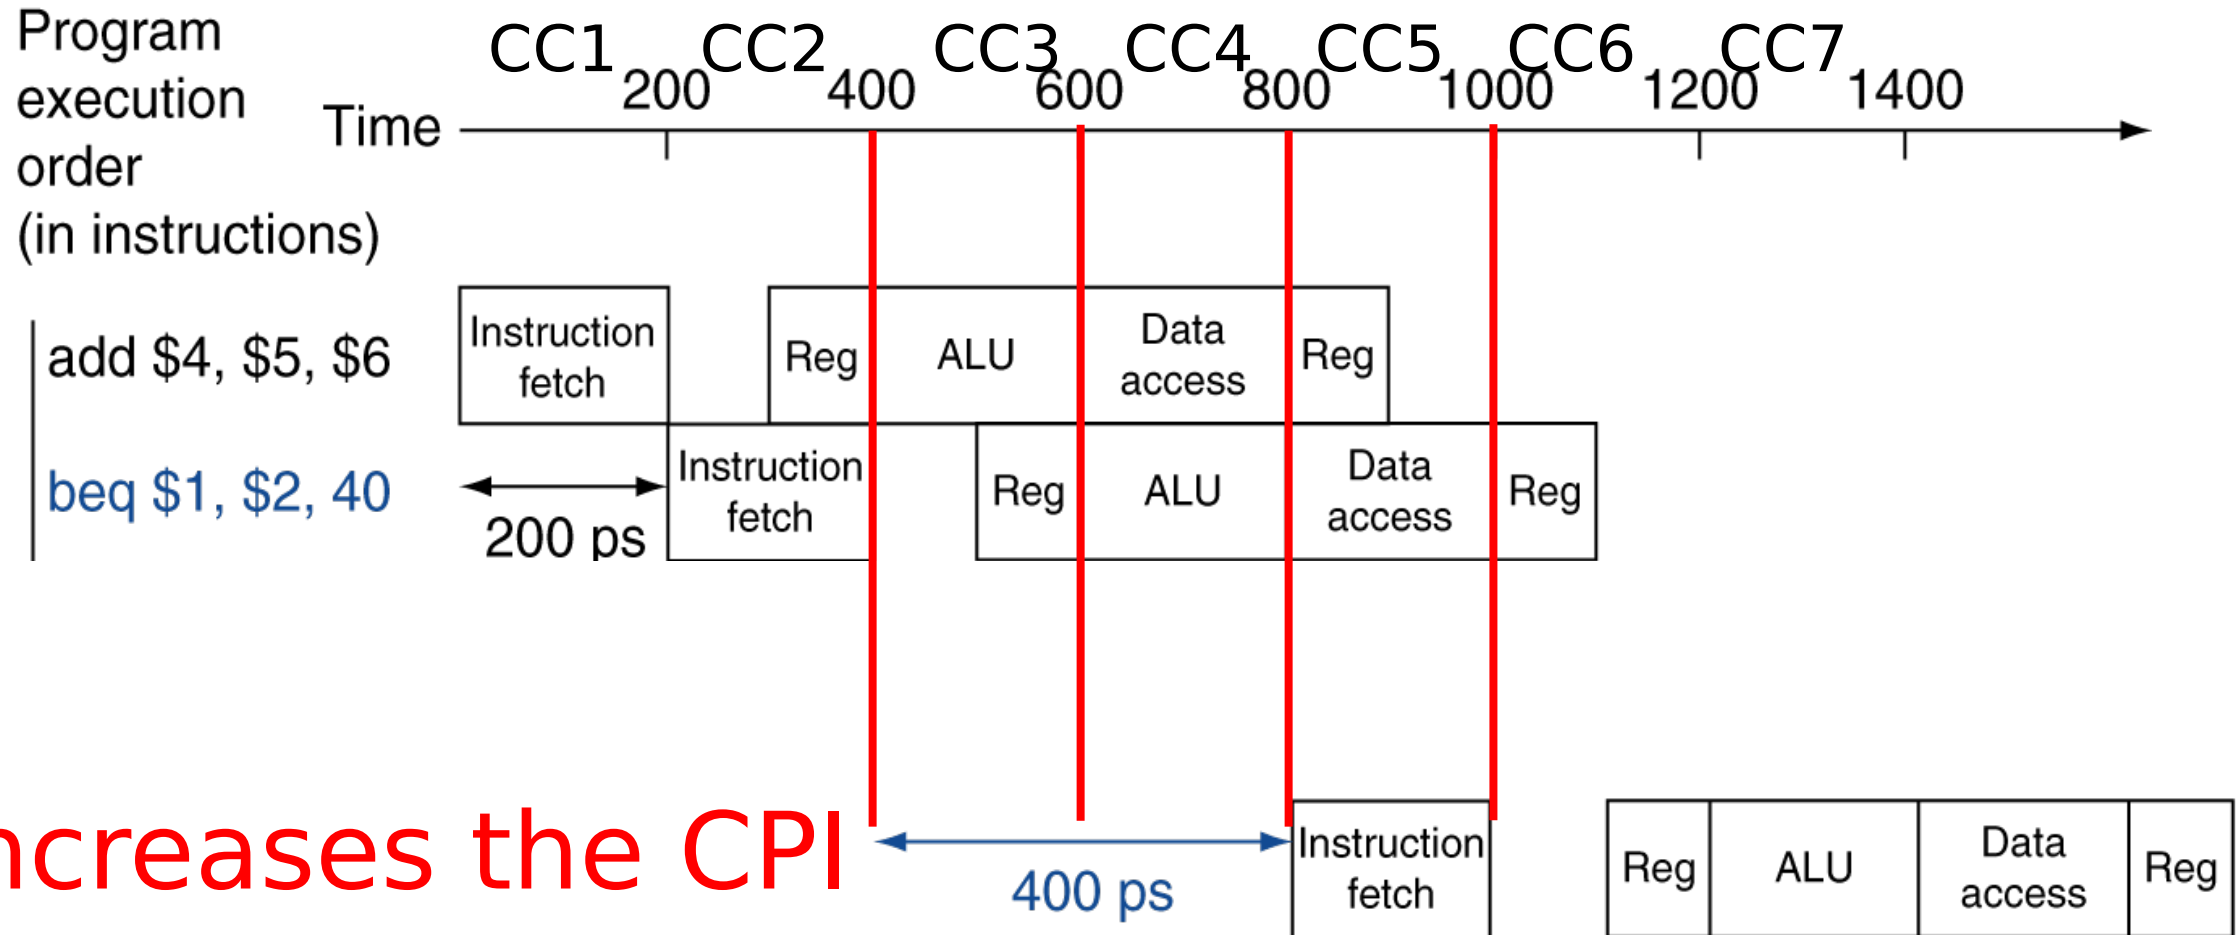
Where are we?

- We have explored a single clock cycle processor design
- We have introduced pipelining and established its performance benefits
- We have explored pipelining hazards and identified solutions



Stall on Branch

- Wait until branch outcome is determined before fetching next instruction

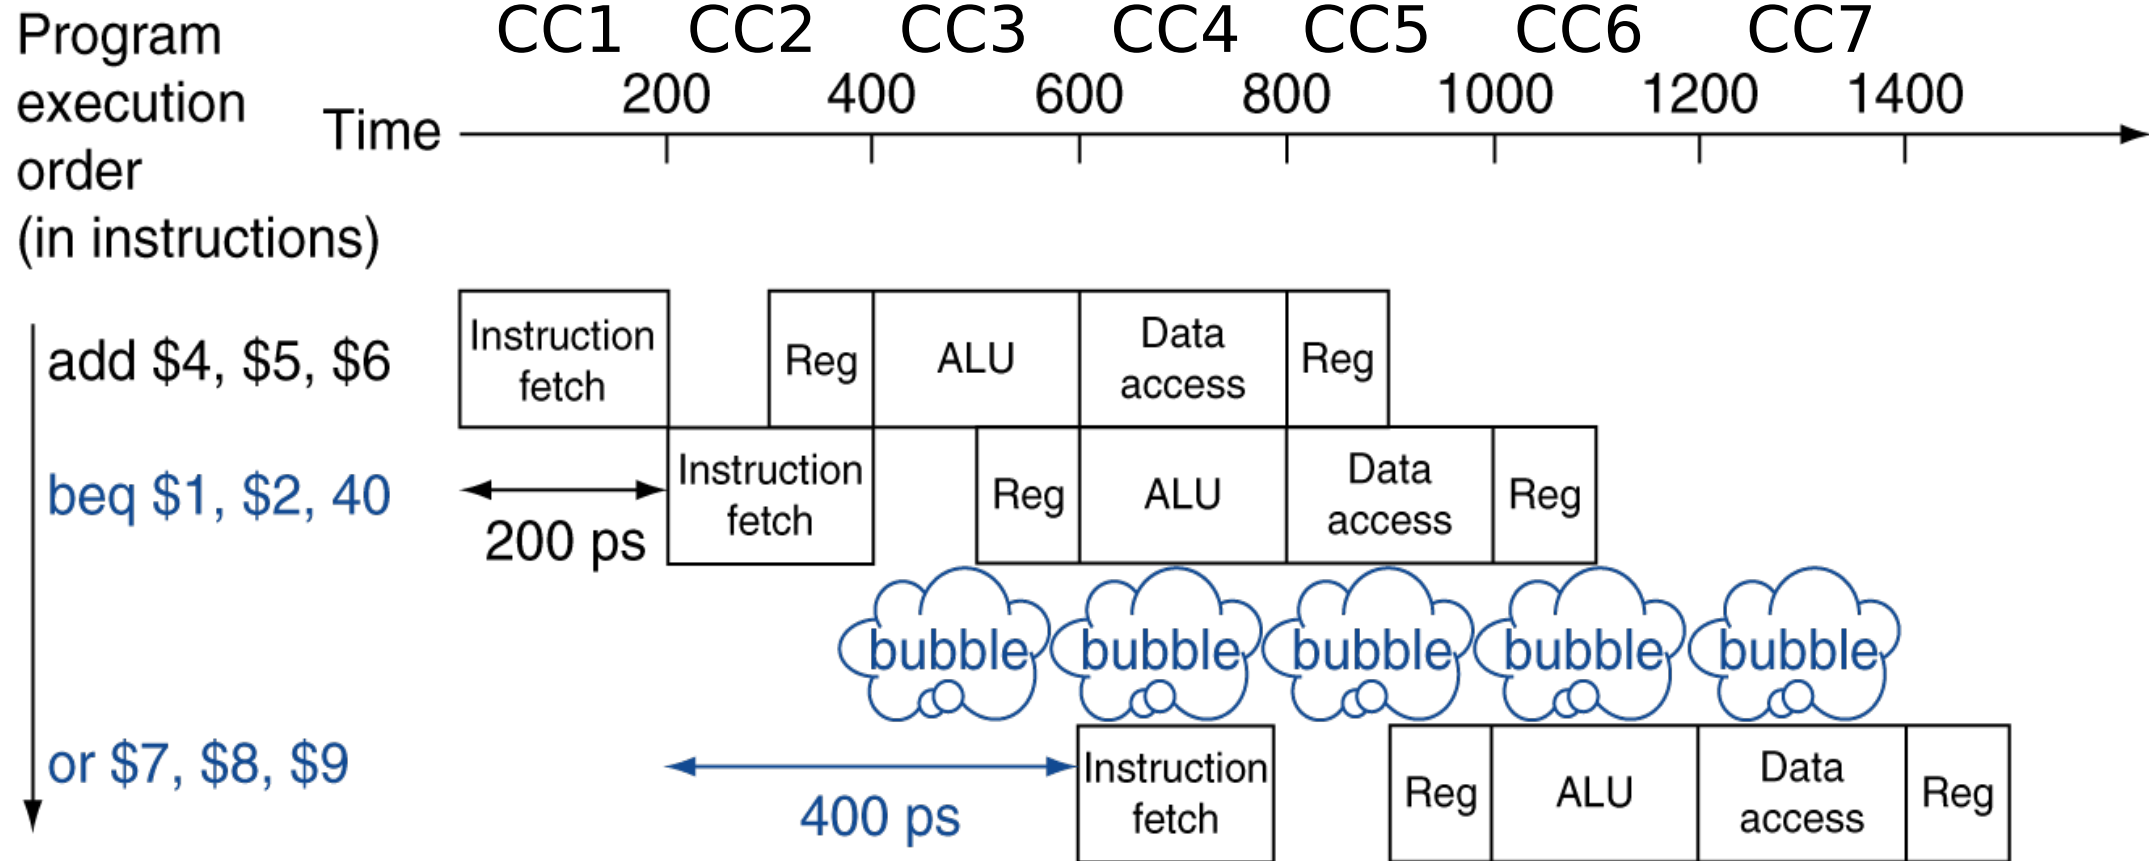


Increases the CPI

Optimized Stall on Branch

Optimized handling using additional hardware

- Decode and decide

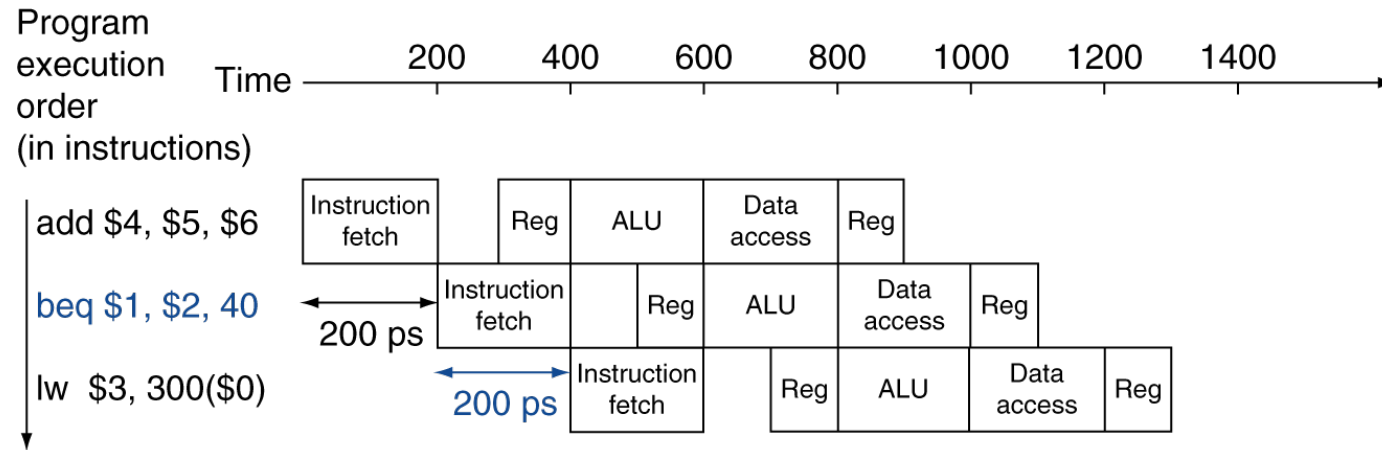


Branch Prediction

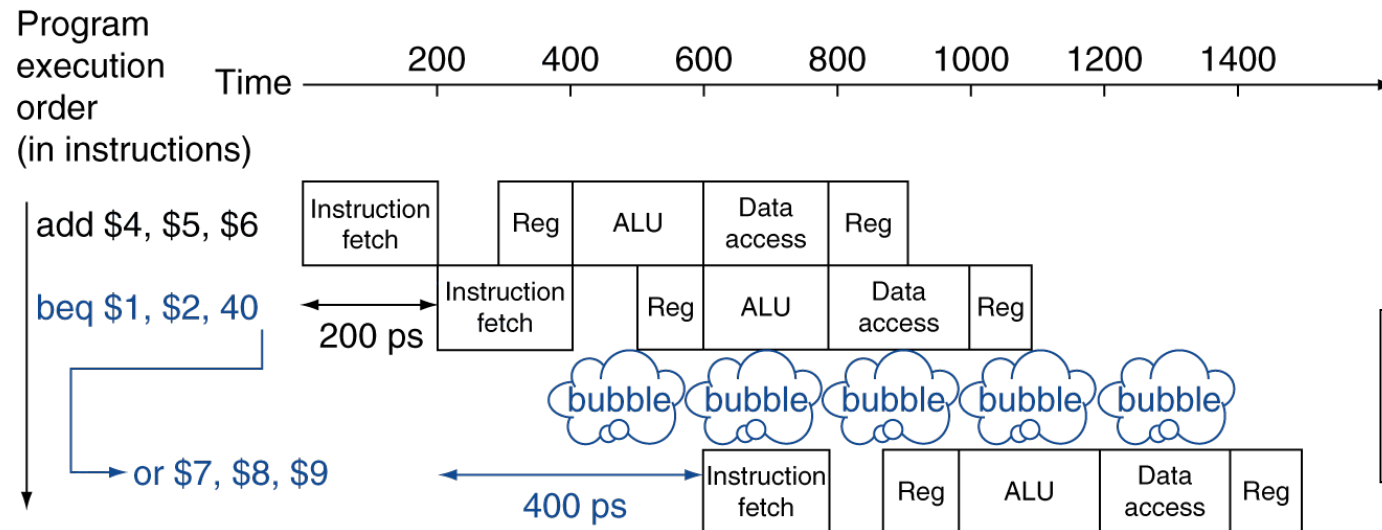
- Usually adopted because Stall penalty becomes unacceptable, **especially for Longer pipelines**
- Predict outcome of branch
 - Only stall if prediction is wrong
 - Accurate prediction → high performance
- In MIPS pipeline
 - Can predict branches not taken



Simple Policy: Predict Branch Not Taken



**Prediction
correct**



**Prediction
incorrect**

**Optimized hardware ->
only one bubble**

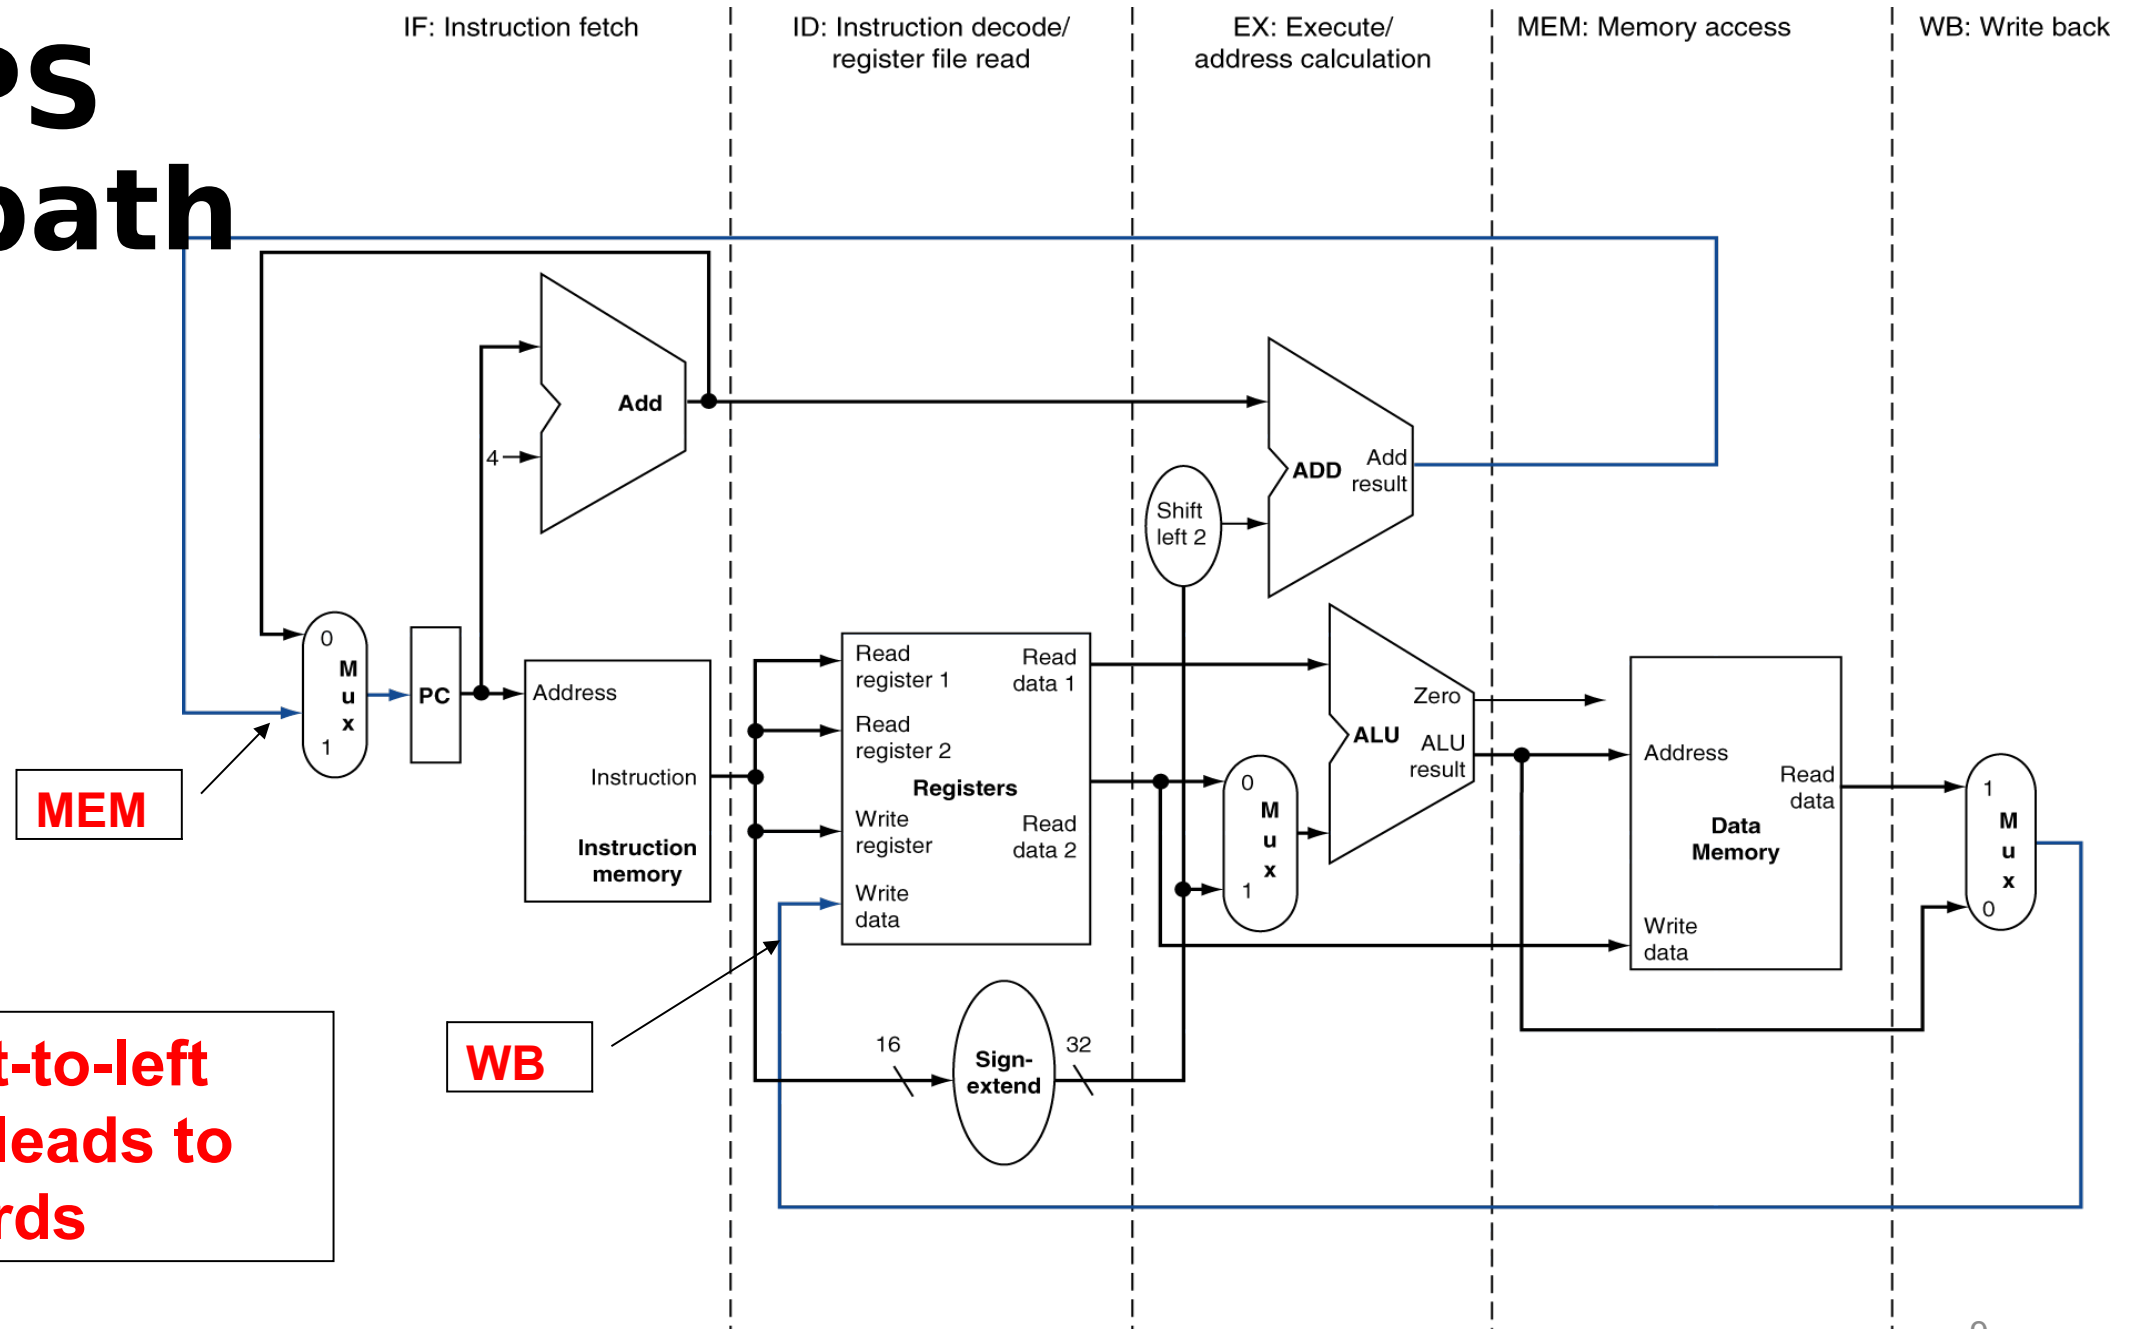
More-Realistic Branch Prediction

- Static branch prediction
 - Based on typical branch behaviour
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- Dynamic branch prediction
 - Hardware measures actual branch behaviour
 - e.g., record recent history of each branch
 - Assume future behaviour will continue the trend

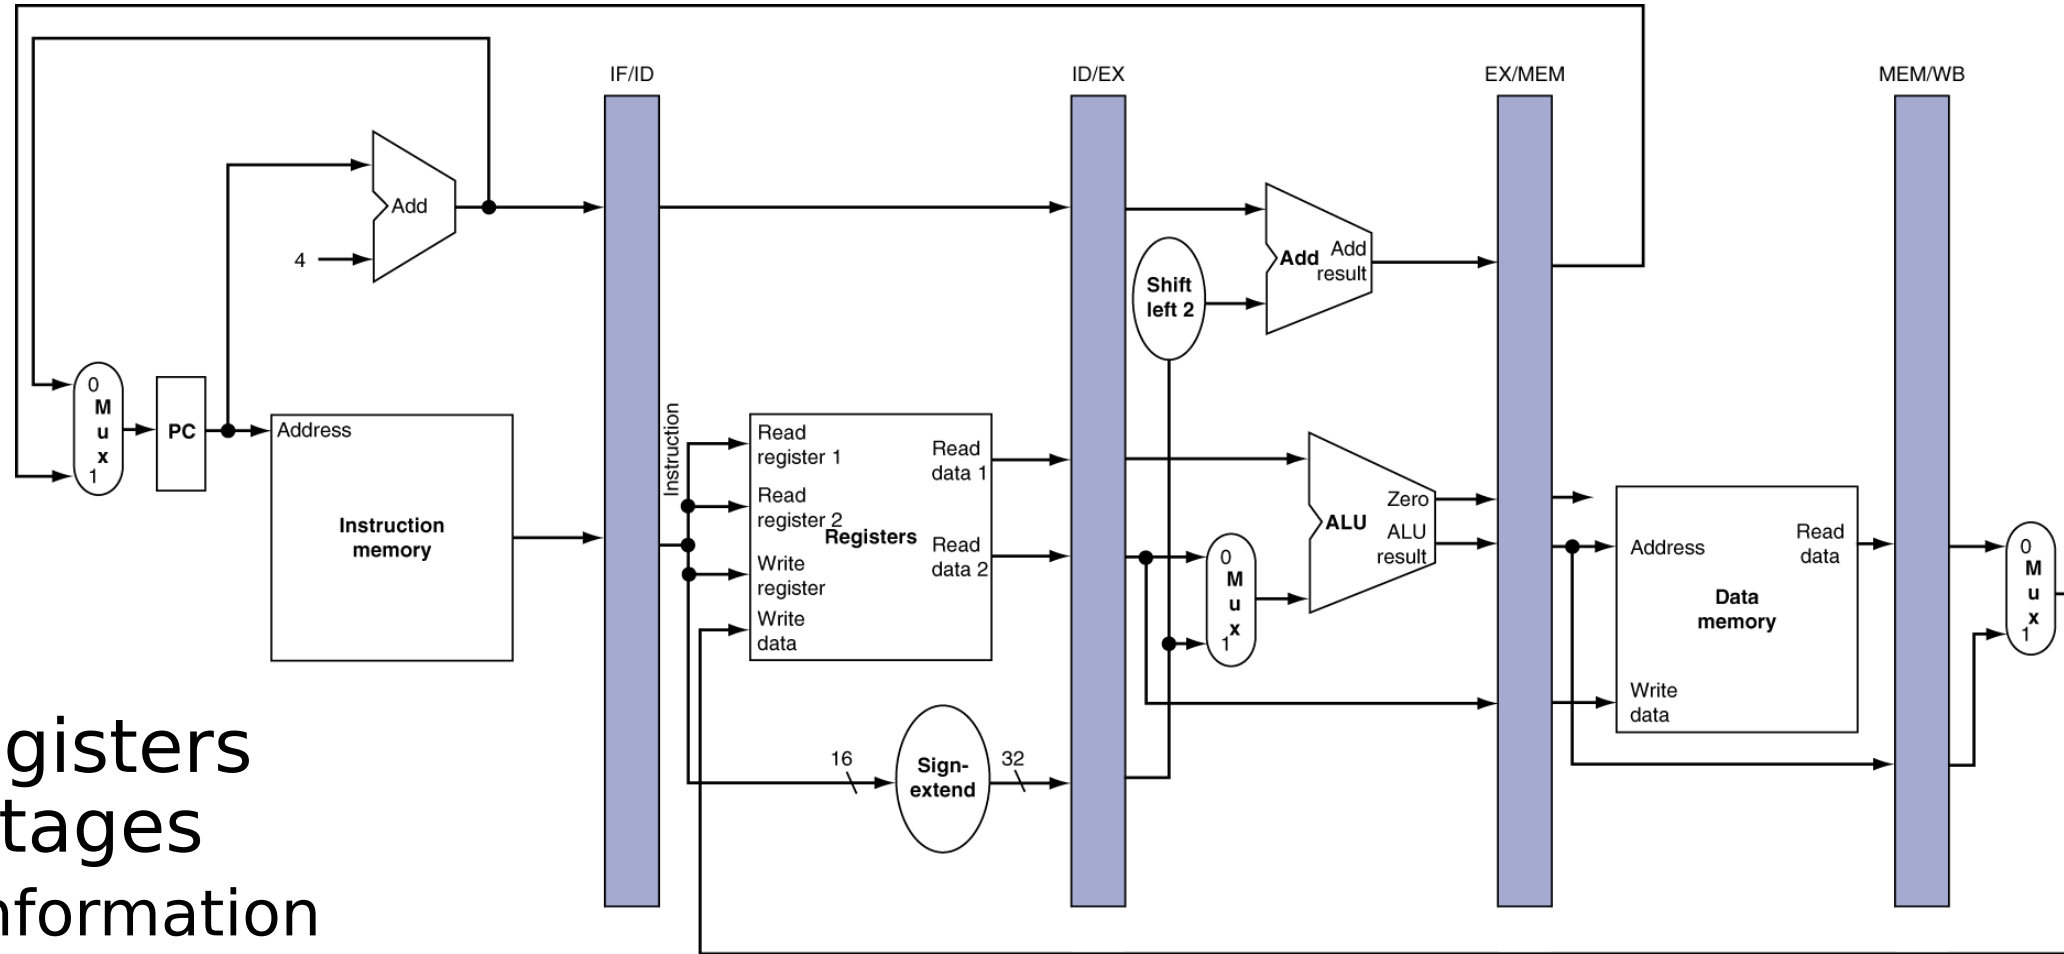
Pipelined Datapath Hardware

- Instructions divided into five stages
 - IF: Instruction fetch
 - ID: Instruction decode and register file read
 - EX: Execution or address calculation
 - MEM: Data memory access
 - WB: Write back
- Dividing instruction into five stages mean five-stage pipeline
 - Up to five instructions could be in execution during any single clock cycle
 - Instruction and data move from left to right through the stages
- Two exceptions to this left-to-right flow of instructions
 - The write back stage
 - Branching

MIPS Datapath



MIPS Pipelined Datapath

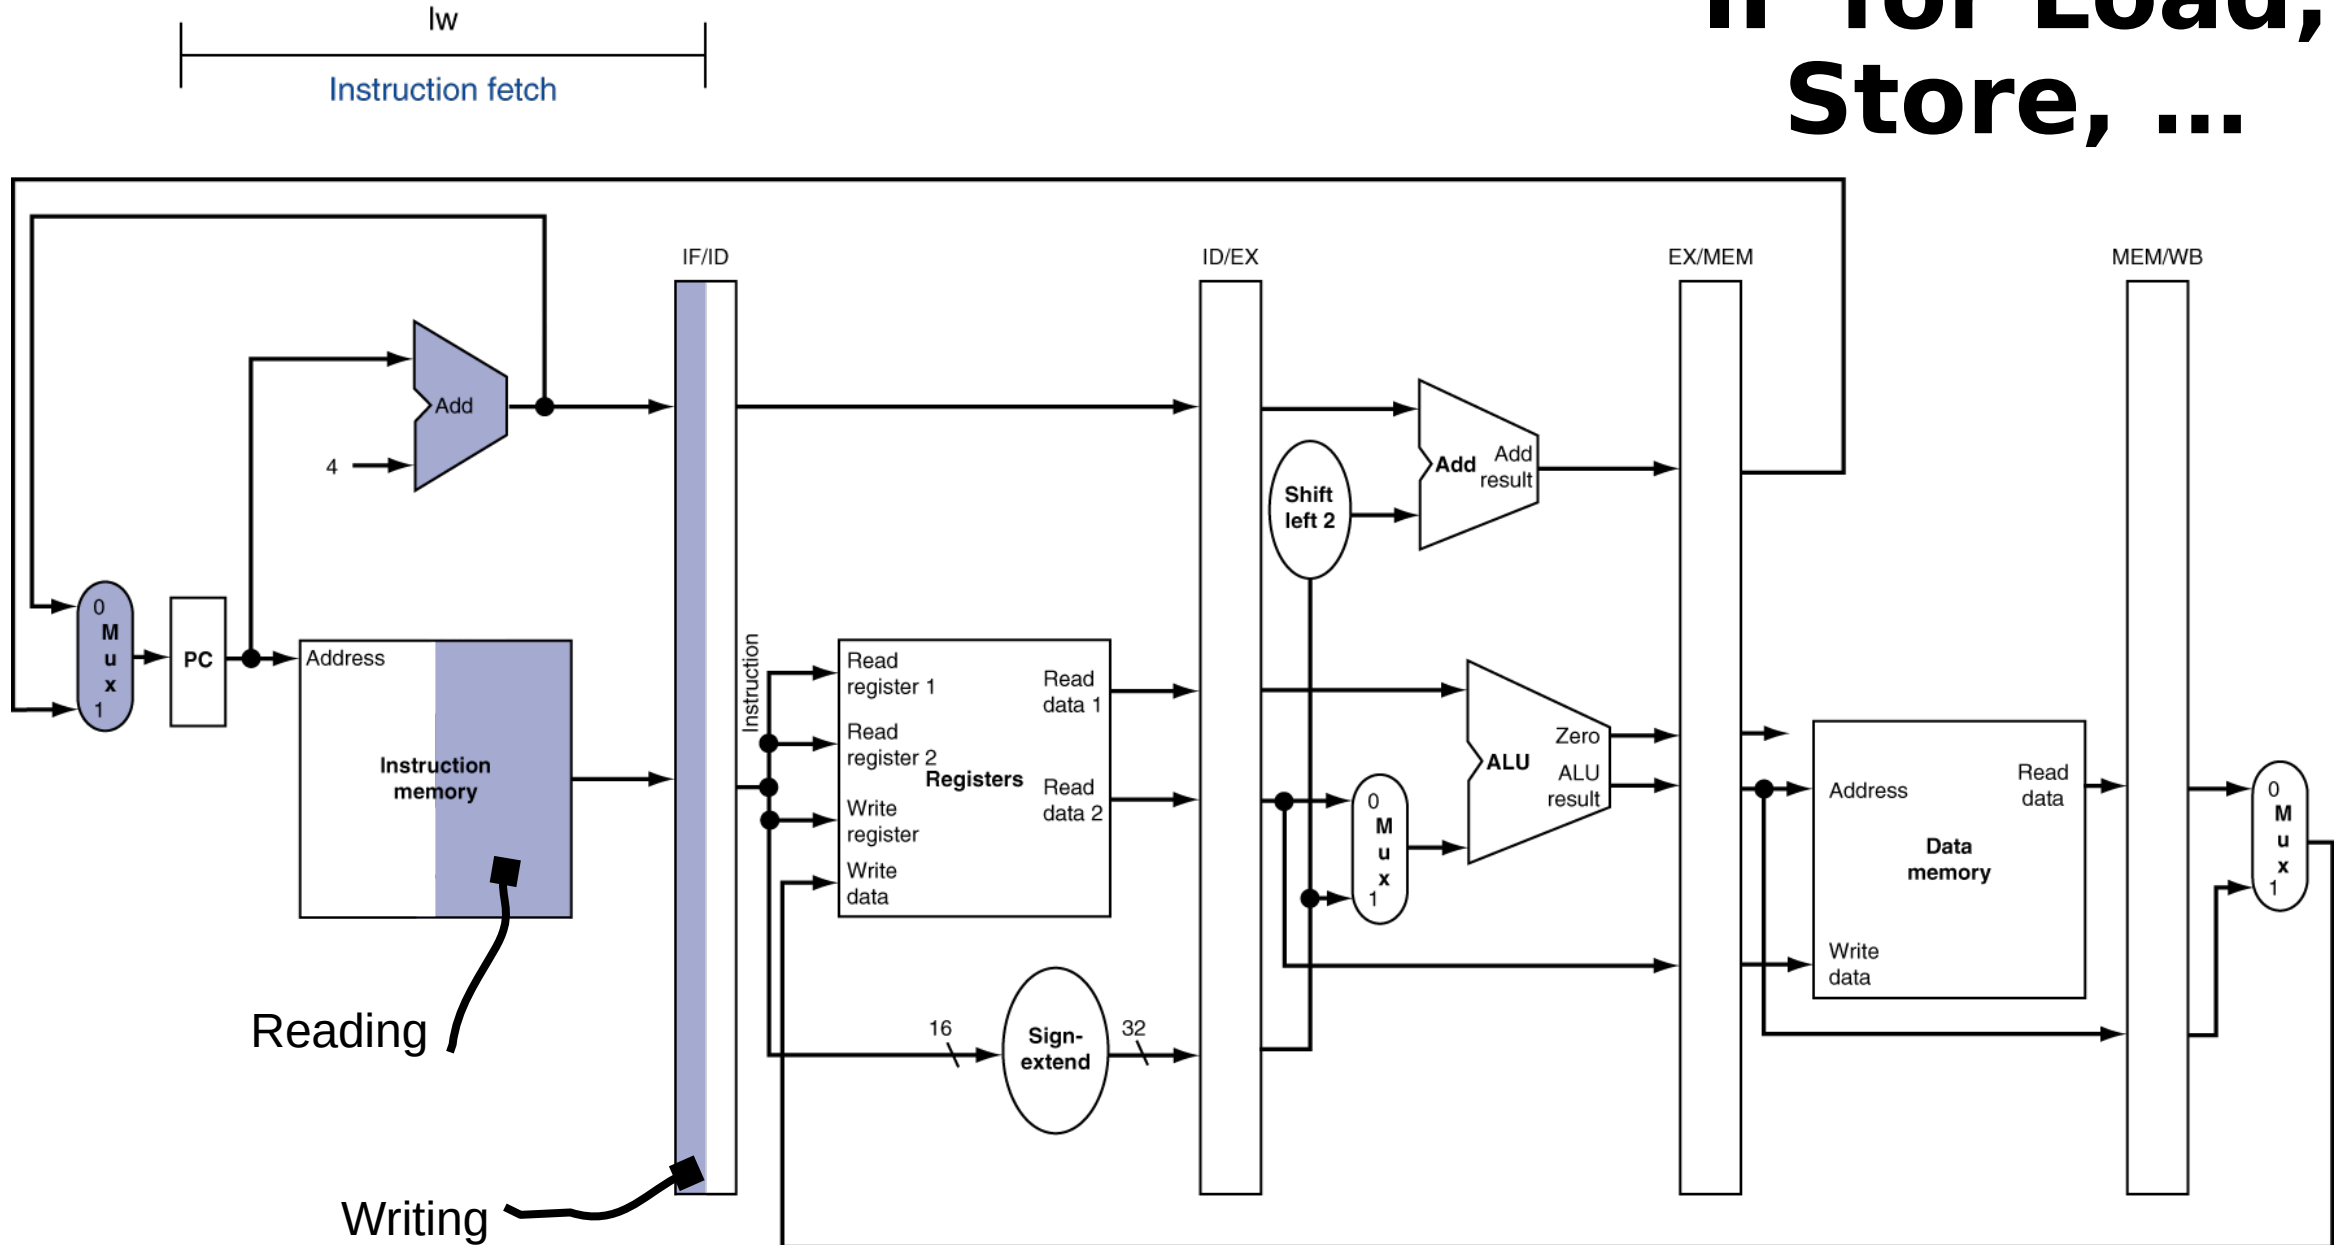


- Pipeline registers between stages
 - To hold information produced in previous cycle

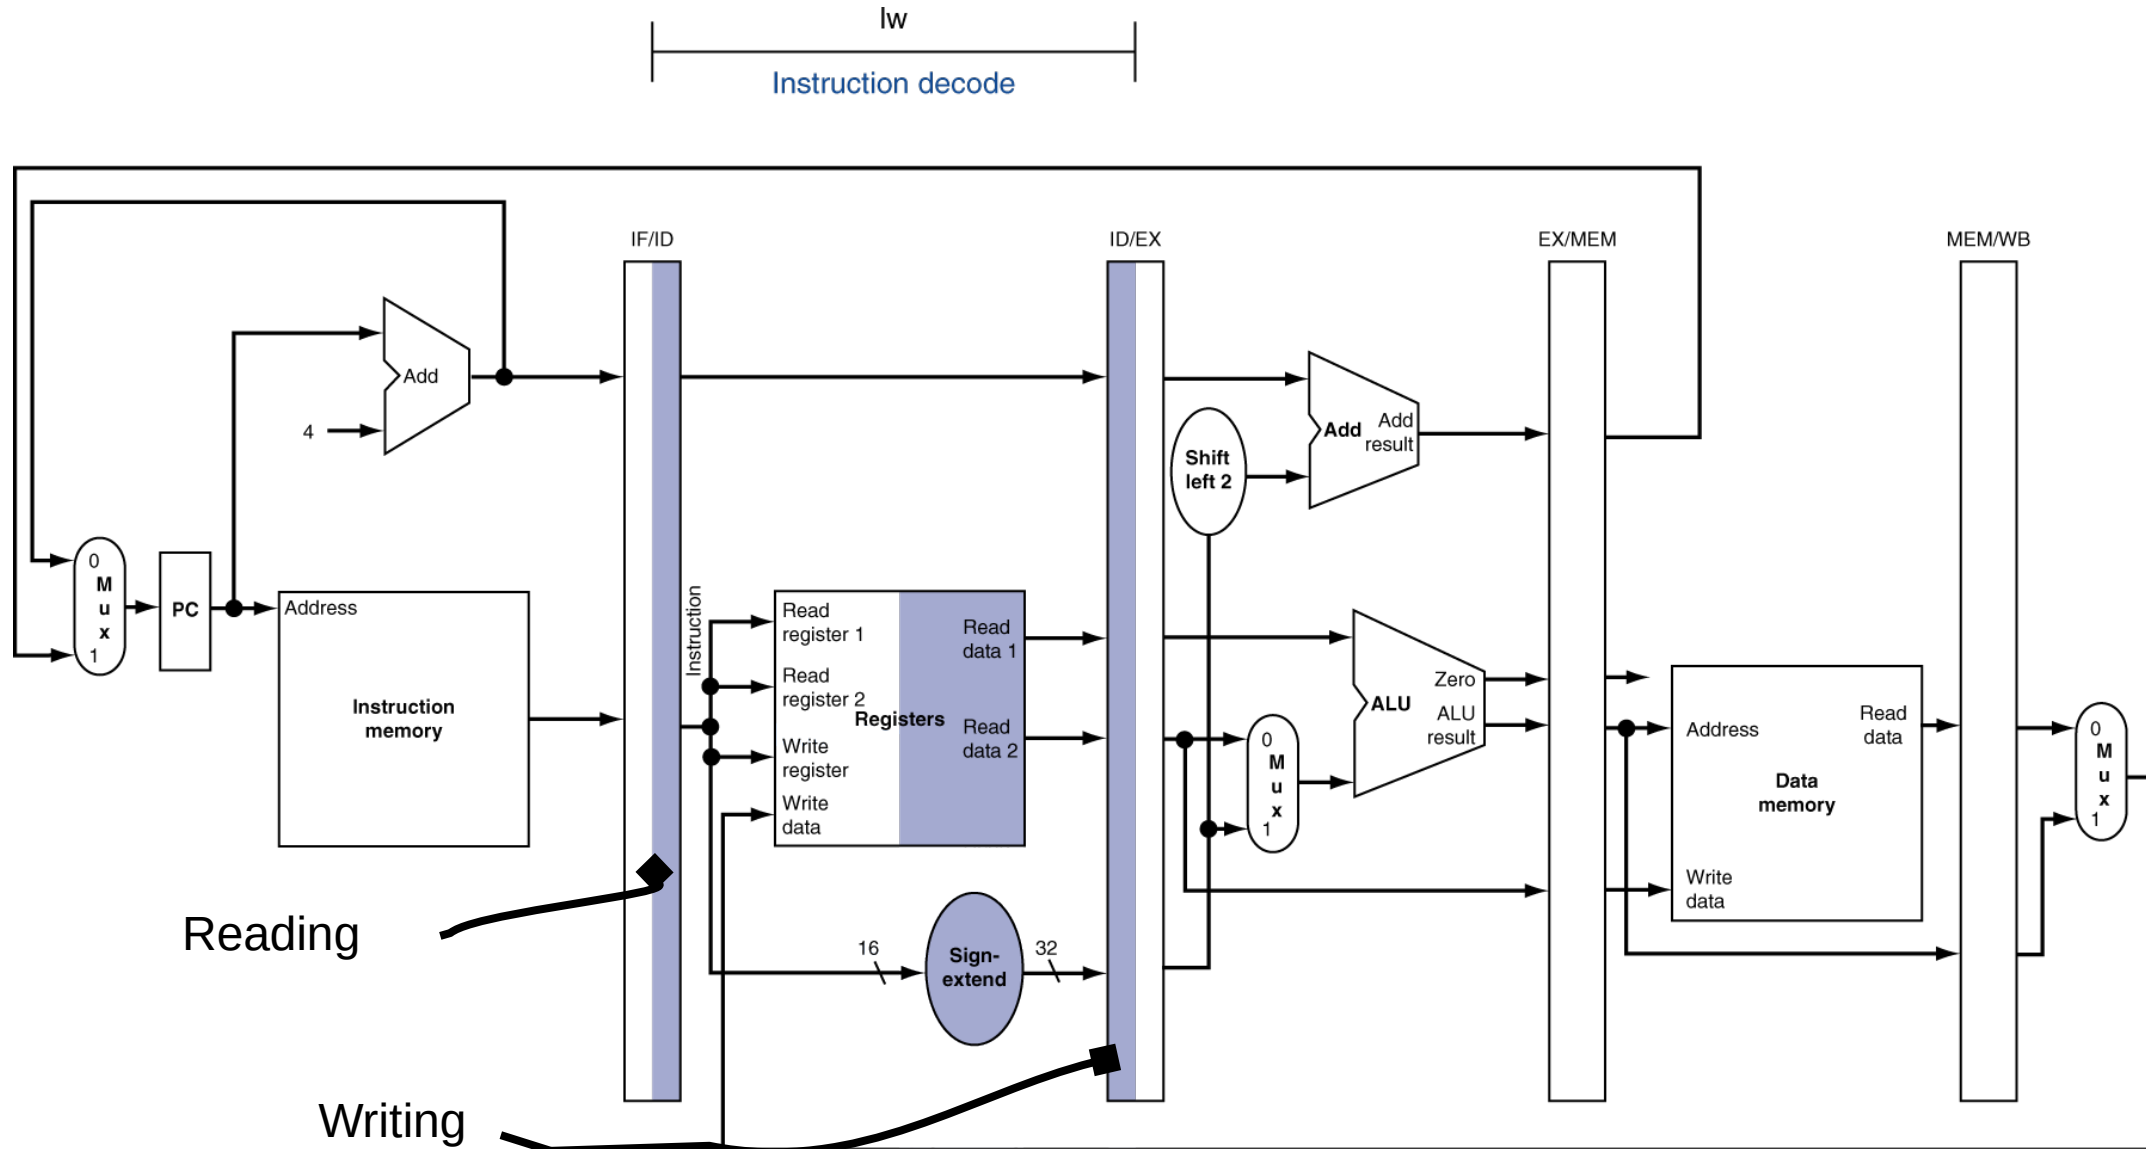
Pipeline Operation Example

- We'll look at “single-clock-cycle” diagrams for load
 - We highlight the right half of register or memory when read operation
 - Highlight left half when they are being written

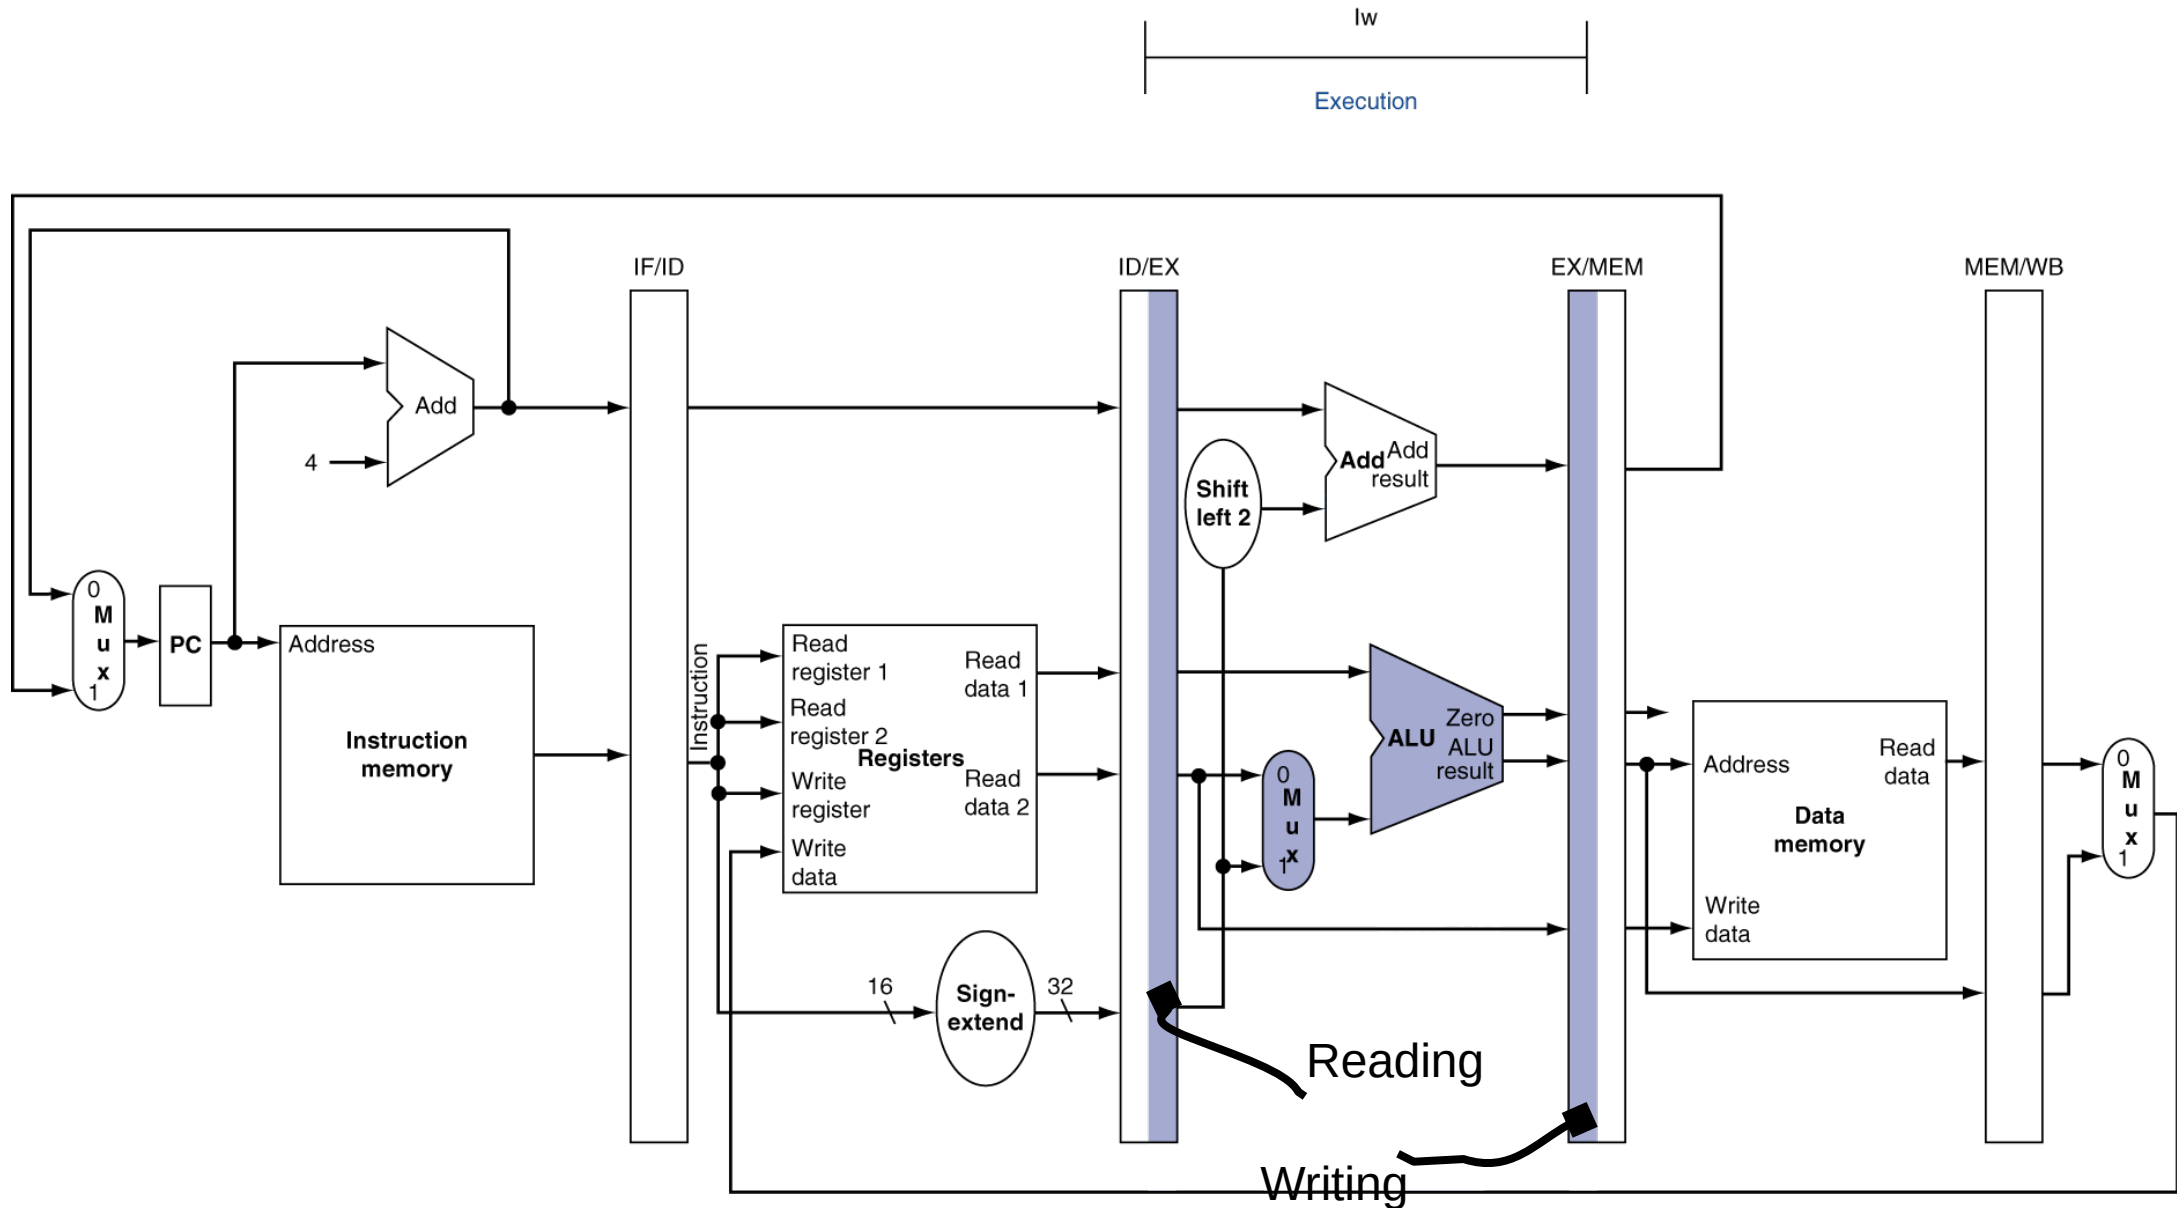
IF for Load, Store, ...



ID for Load, Store, ...

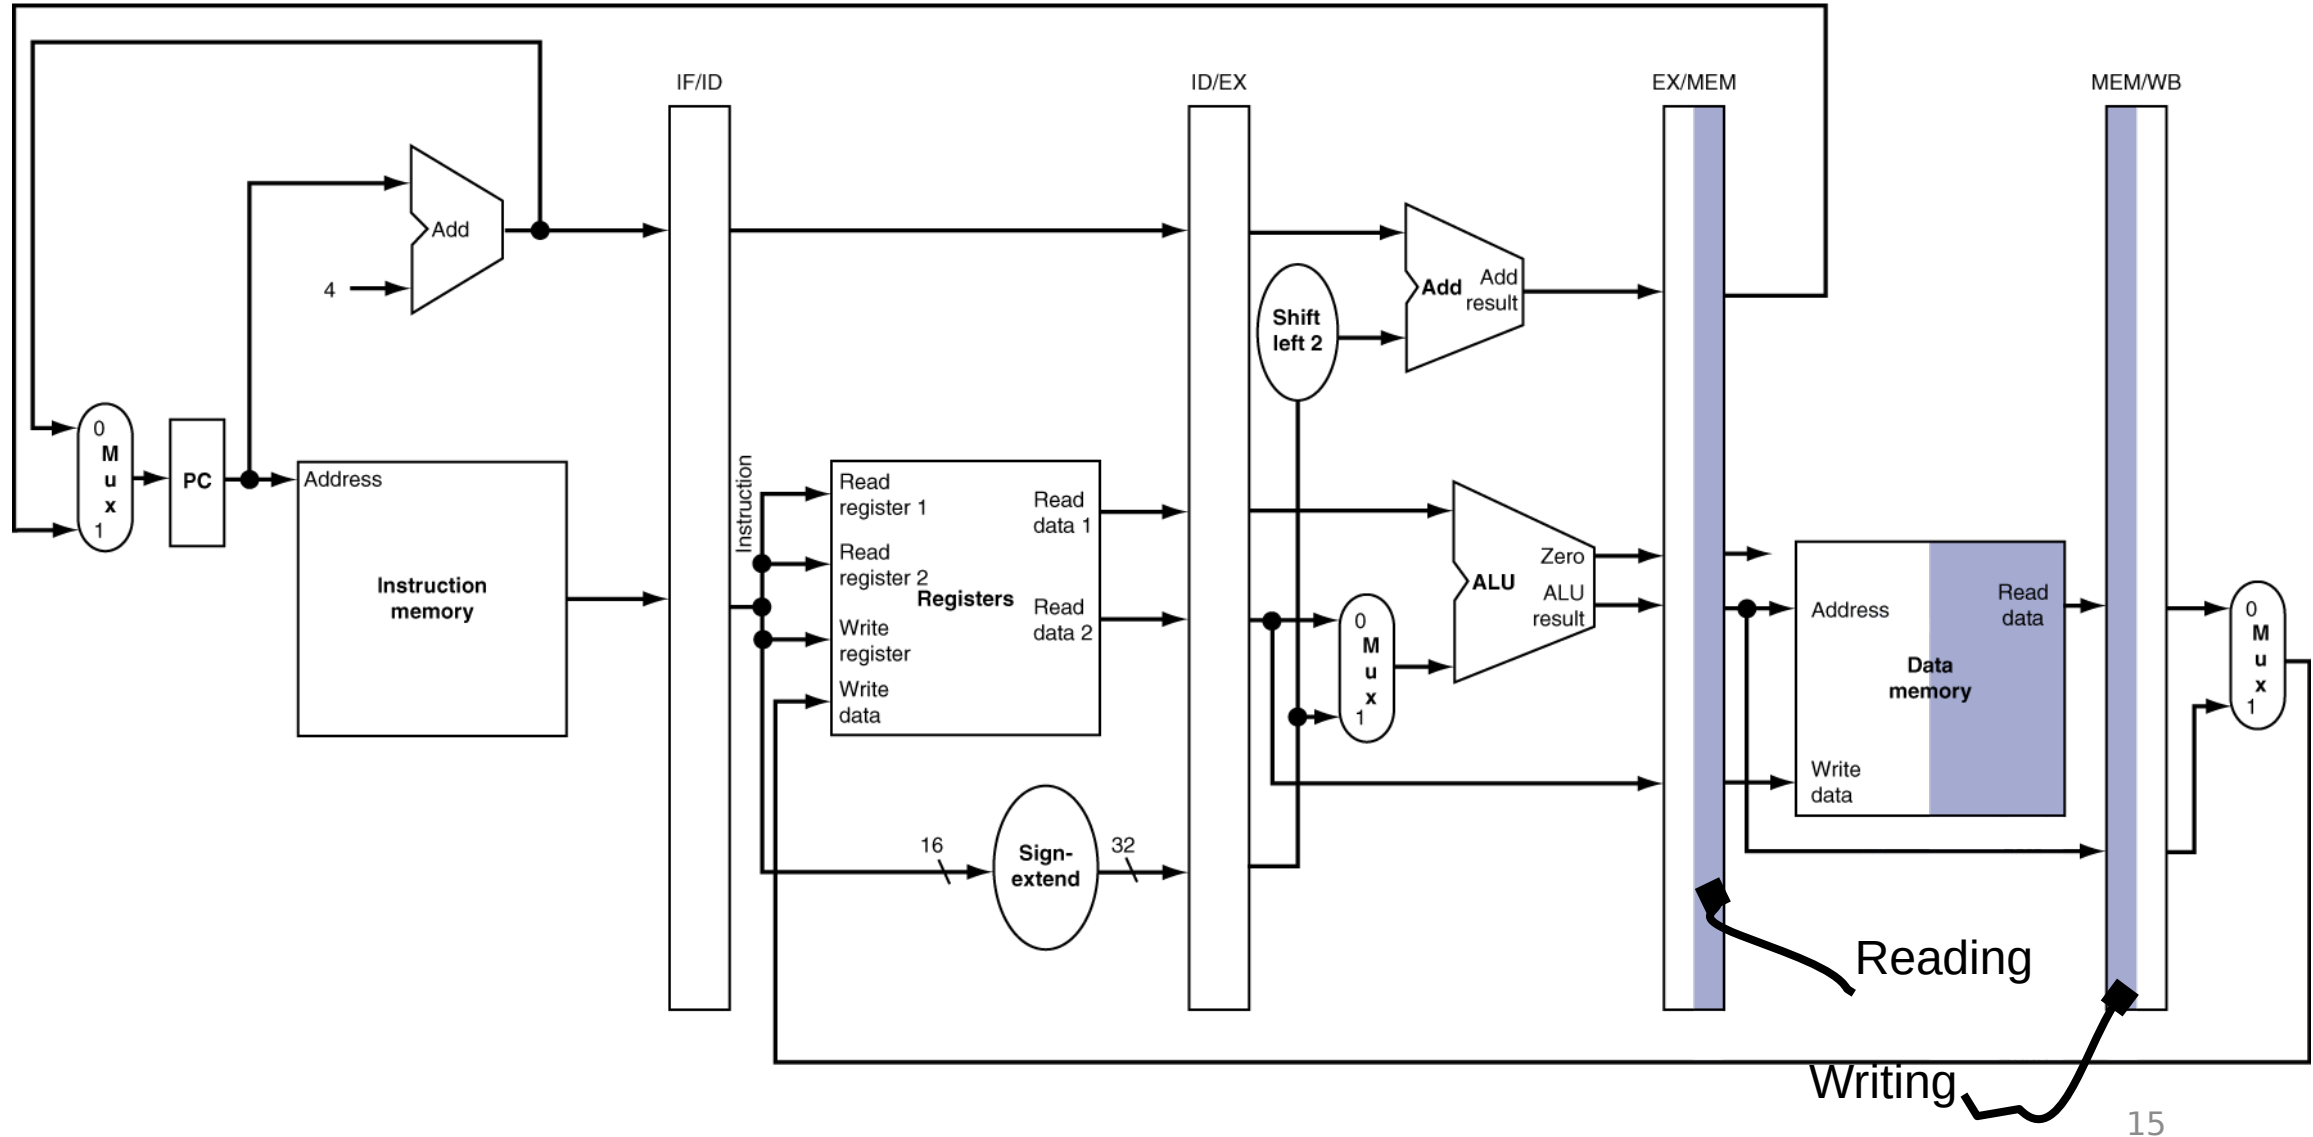


EX for Load



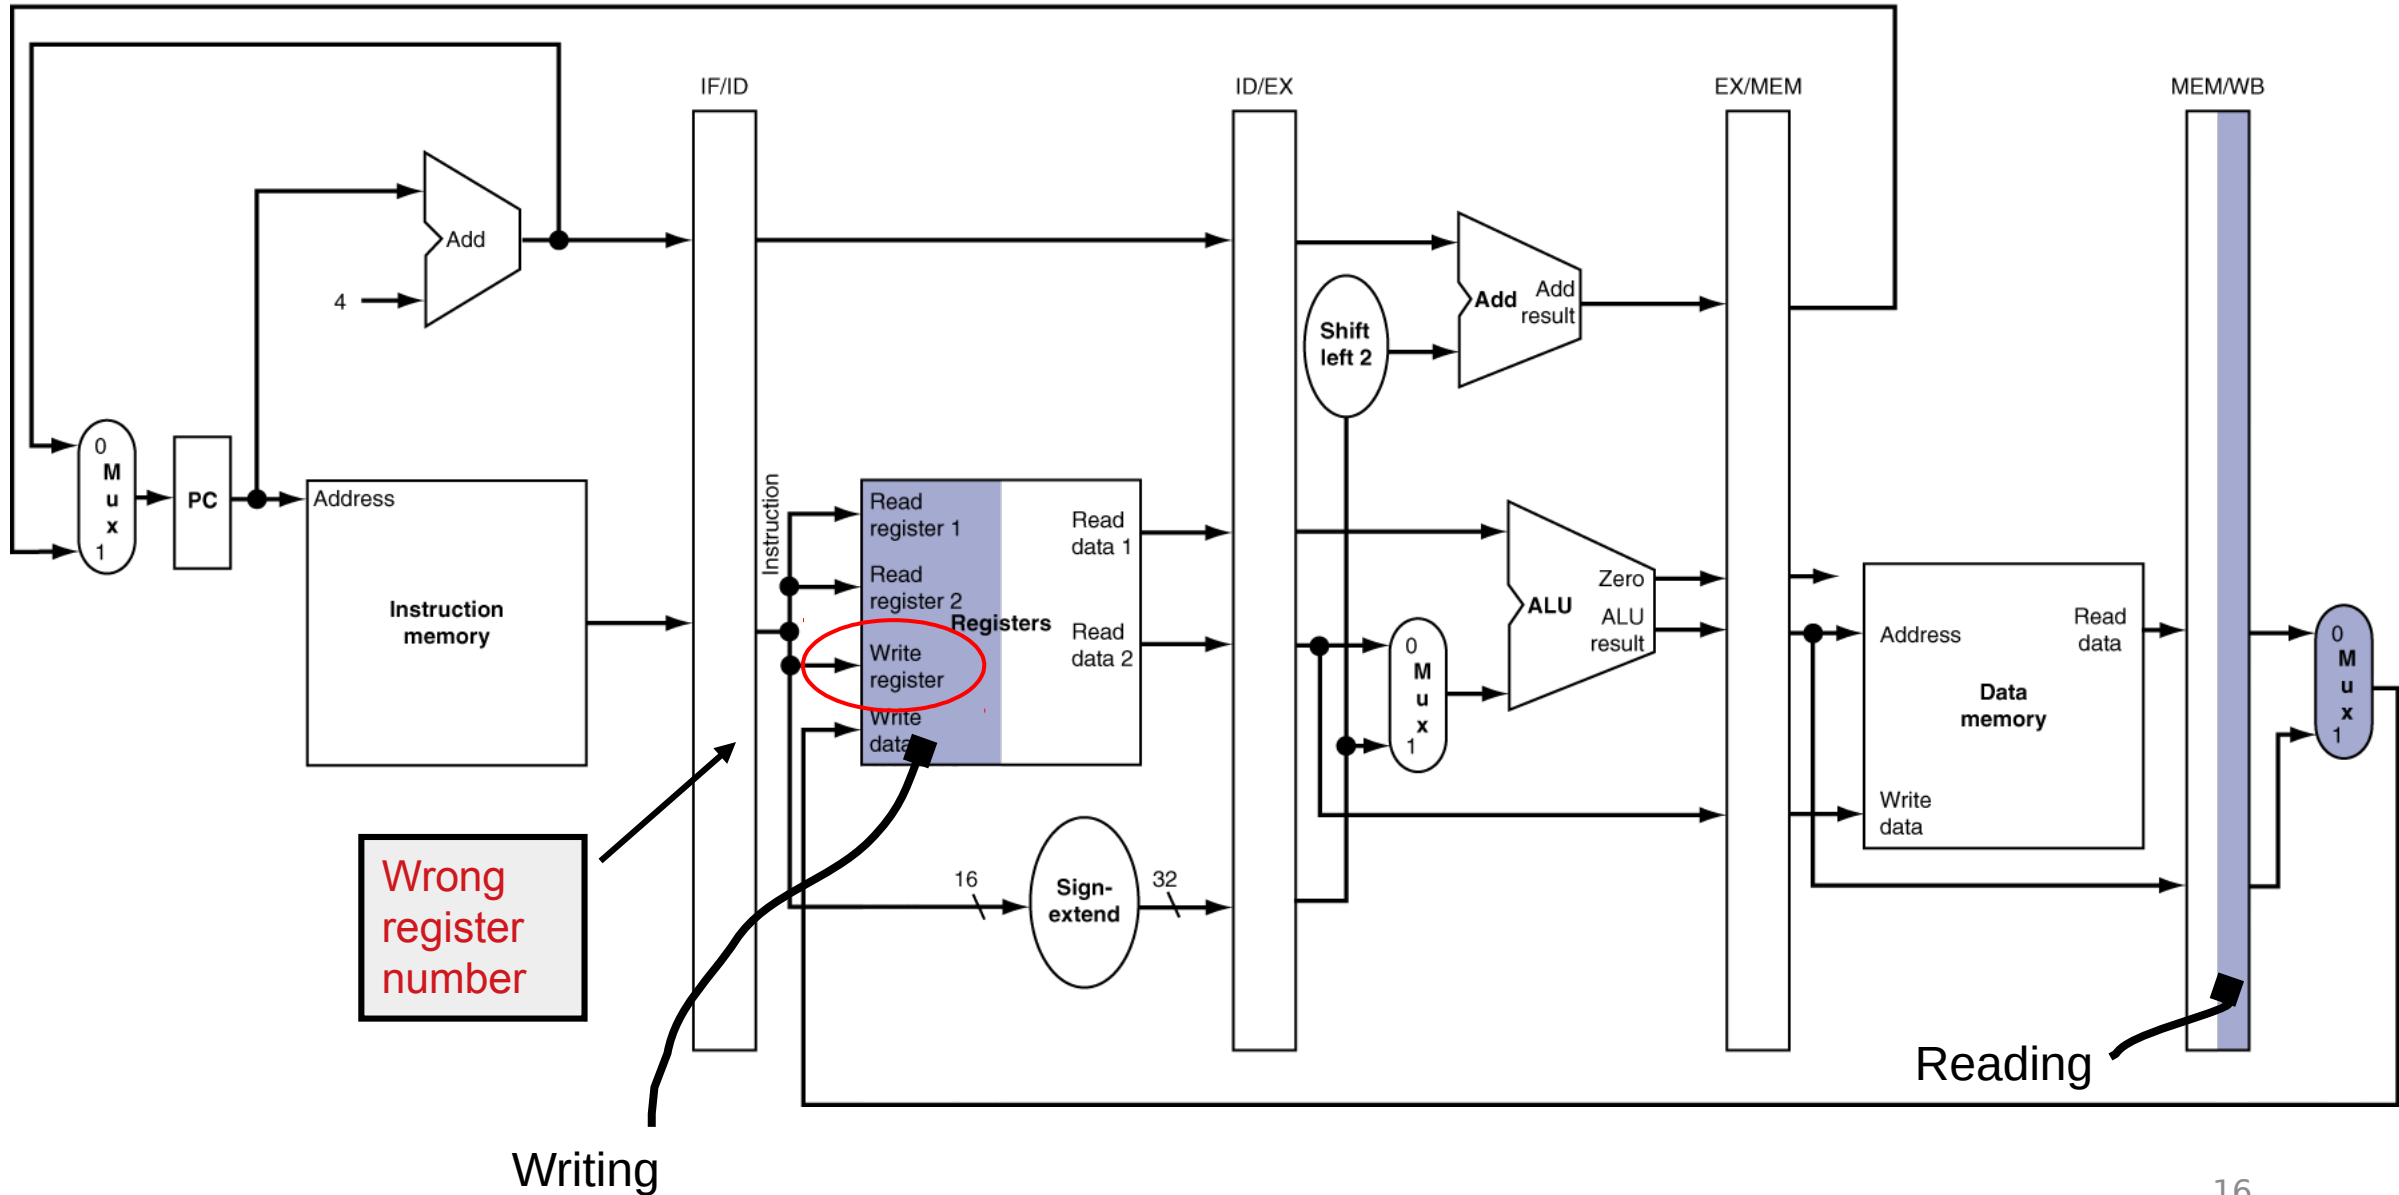
MEM for Load

lw
Memory




WB for Load

lw
Write back



Pipelined Control

- Control signals derived from instruction
- Control lines start with execution stage
- Relevant control should propagate with the instruction → Extended stage registers



Pipelined Control Design `

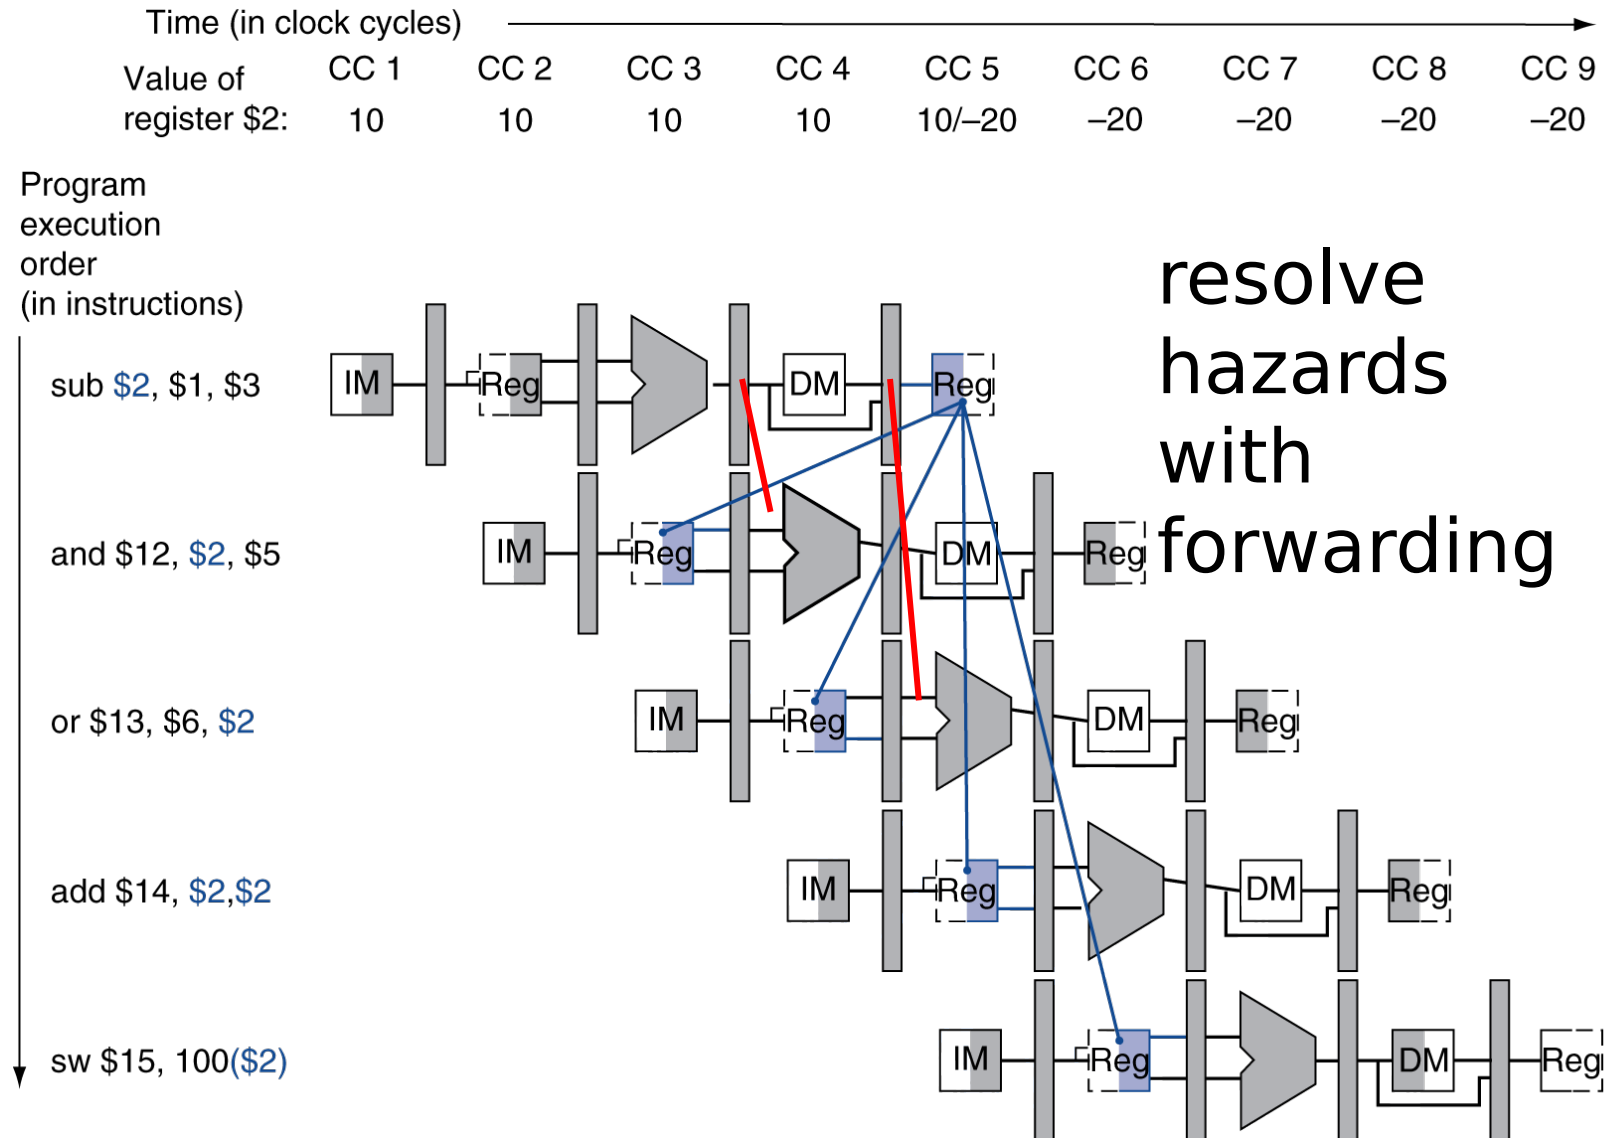
Data Hazards in ALU Instructions

- Consider this sequence with many dependences:

sub \$2, \$1, \$3
 and \$12, \$2, \$5
 or \$13, \$6, \$2
 add \$14, \$2, \$2
 sw \$15, 100(\$2)

Assuming \$2 has the value


10 before the subtraction
 -20 after the subtraction



Forwarding Hardware

- Forward implies reading ALU operands from a source other than the register file registers
 - **Multiplexers** are needed at the ALU I/P
 - A **forwarding unit** determines the multiplexor control signal

Forwarding Paths



b. With forwarding

Forwarding Unit Output

Mux control	Source	Explanation
ForwardA = 00	ID/EX	The first ALU operand comes from the register file.
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result.
ForwardA = 01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result.
ForwardB = 00	ID/EX	The second ALU operand comes from the register file.
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result.
ForwardB = 01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result.

Reading

Section 4.6