


Where are we?

1. Where can a block be placed? (Q1)
2. How is a block found? (Q2)
3. What block is replaced on a miss? (Q3)
4. How are writes handled? (Q4)

Memory Hierarchy



Direct mapped cache memory


Cache Performance

Improving Cache Performance

Reduce miss rate (??)

Reduce miss (time) penalty (??)

Speed hit access time!



Example

- We have **4-block** caches
 - Direct mapped
 - **Block access sequence:** 0, 8, 0, 6, 8

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|---|--------|---|
| | | | 0 | 1 | 2 | 3 |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[8] | | | |
| 0 | 0 | miss | Mem[0] | | | |
| 6 | 2 | miss | Mem[0] | | Mem[6] | |
| 8 | 0 | miss | Mem[8] | | Mem[6] | |

Reducing Miss Rate

Flexible block location to reduce the competition for cache blocks → **Associative Caches**

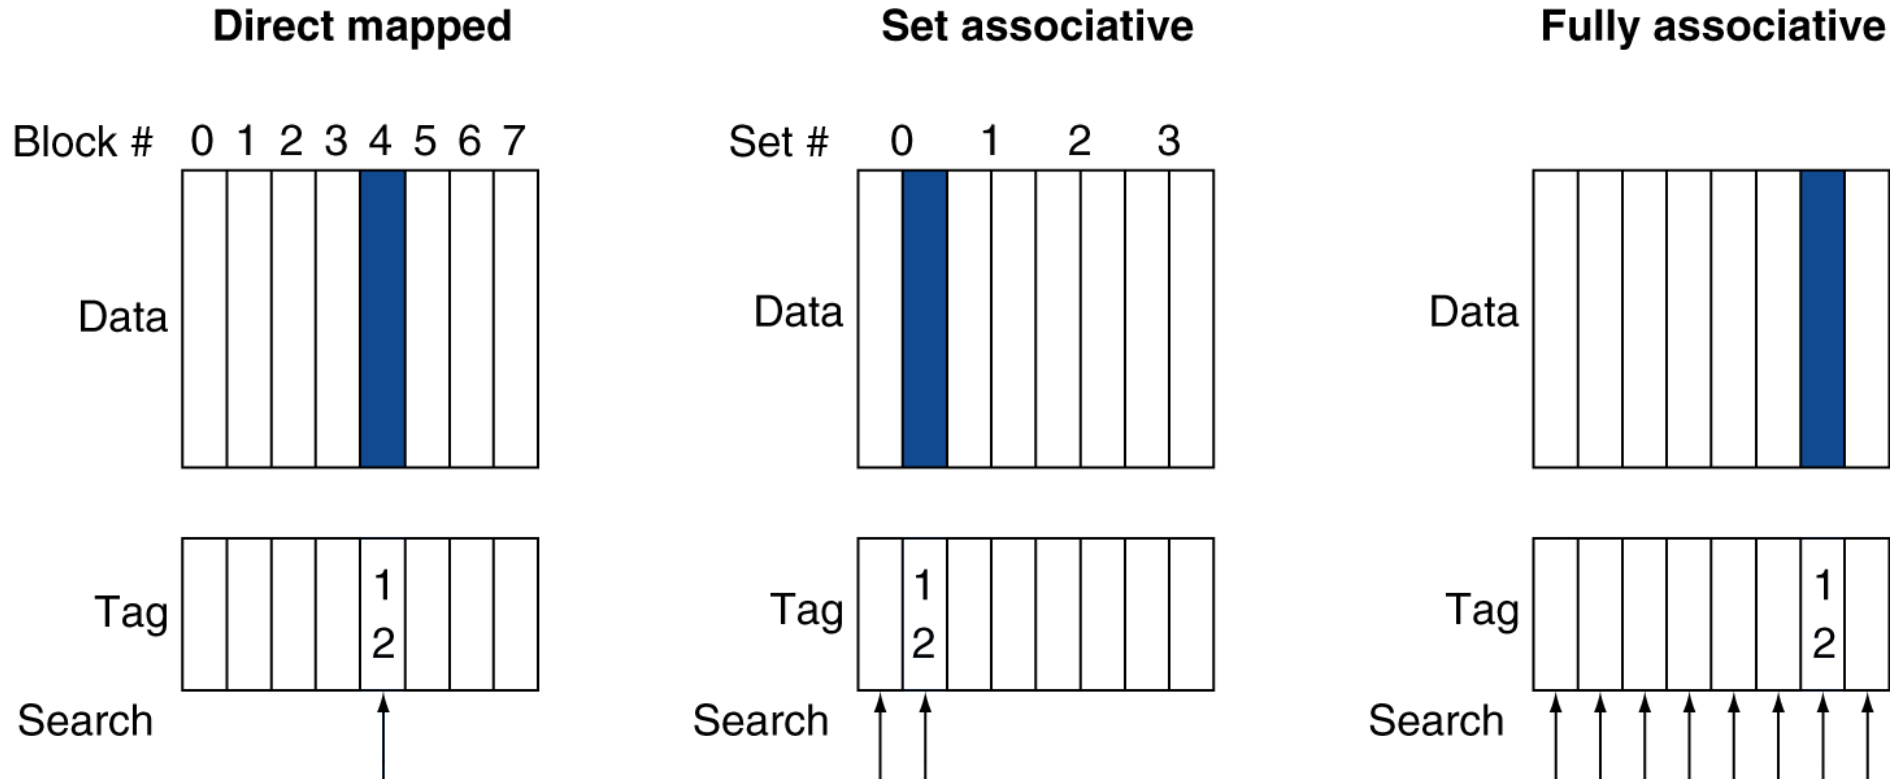
- ***Fully associative caches***

- Allow a given block to go in any cache entry (less competition)
- Increased search time
 - Search all entries at once to reduce hit time

- ***n-way set-associative caches***

- Each set contains n entries
- Block number determines which set
 - (Block number) modulo (#Sets in cache)
- Search all entries in a given set at once
 - n comparators (less expensive)

Associative Cache Example



***Block position =
(Block #) modulo
(# of blocks in the cache)***

***Set position =
(Block #) modulo
(# of sets in the cache)***

Spectrum of Associativity

- For a cache with 8 entries

One-way set associative (direct mapped)

| Block | Tag | Data |
|-------|-----|------|
| 0 | | |
| 1 | | |
| 2 | | |
| 3 | | |
| 4 | | |
| 5 | | |
| 6 | | |
| 7 | | |

Two-way set associative

| Set | Tag | Data | Tag | Data |
|-----|-----|------|-----|------|
| 0 | | | | |
| 1 | | | | |
| 2 | | | | |
| 3 | | | | |

Four-way set associative

| Set | Tag | Data | Tag | Data | Tag | Data | Tag | Data |
|-----|-----|------|-----|------|-----|------|-----|------|
| 0 | | | | | | | | |
| 1 | | | | | | | | |

Eight-way set associative (fully associative)

| Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data |
|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|
| | | | | | | | | | | | | | | | |

Associativity Example

- 2-way set associative

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|--------|-------|--|
| | | | Set 0 | | Set 1 | |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[0] | Mem[8] | | |
| 0 | 0 | hit | Mem[0] | Mem[8] | | |
| 6 | 0 | miss | Mem[0] | Mem[6] | | |
| 8 | 0 | miss | Mem[8] | Mem[6] | | |

Associativity Example


Fully associative

| Block address | | Hit/miss | Cache content after access | | | |
|---------------|--|----------|----------------------------|---------------|---------------|--|
| 0 | | miss | Mem[0] | | | |
| 8 | | miss | Mem[0] | Mem[8] | | |
| 0 | | hit | Mem[0] | Mem[8] | | |
| 6 | | miss | Mem[0] | Mem[8] | Mem[6] | |
| 8 | | hit | Mem[0] | Mem[8] | Mem[6] | |

How Much Associativity?

- **Increased associativity decreases miss rate**
 - But with diminishing returns
- Simulation of a system with 64KB D-cache, 16-word blocks, SPEC2000
 - 1-way: 10.3%
 - 2-way: 8.6%
 - 4-way: 8.3%
 - 8-way: 8.1%
- Performance cost tradeoff
 - Better performance requires additional hardware for speeding hit time

Set Associative Cache Organization




Replacement Policy

- Direct mapped: no choice
- Set associative
 - Prefer non-valid entry, if there is one
 - Otherwise, choose among entries in the set
- Least-recently used (LRU)
 - Choose the one unused for the longest time
 - Simple for 2-way, manageable for 4-way, too hard beyond that
- Random
 - Gives approximately the same performance as LRU for high associativity

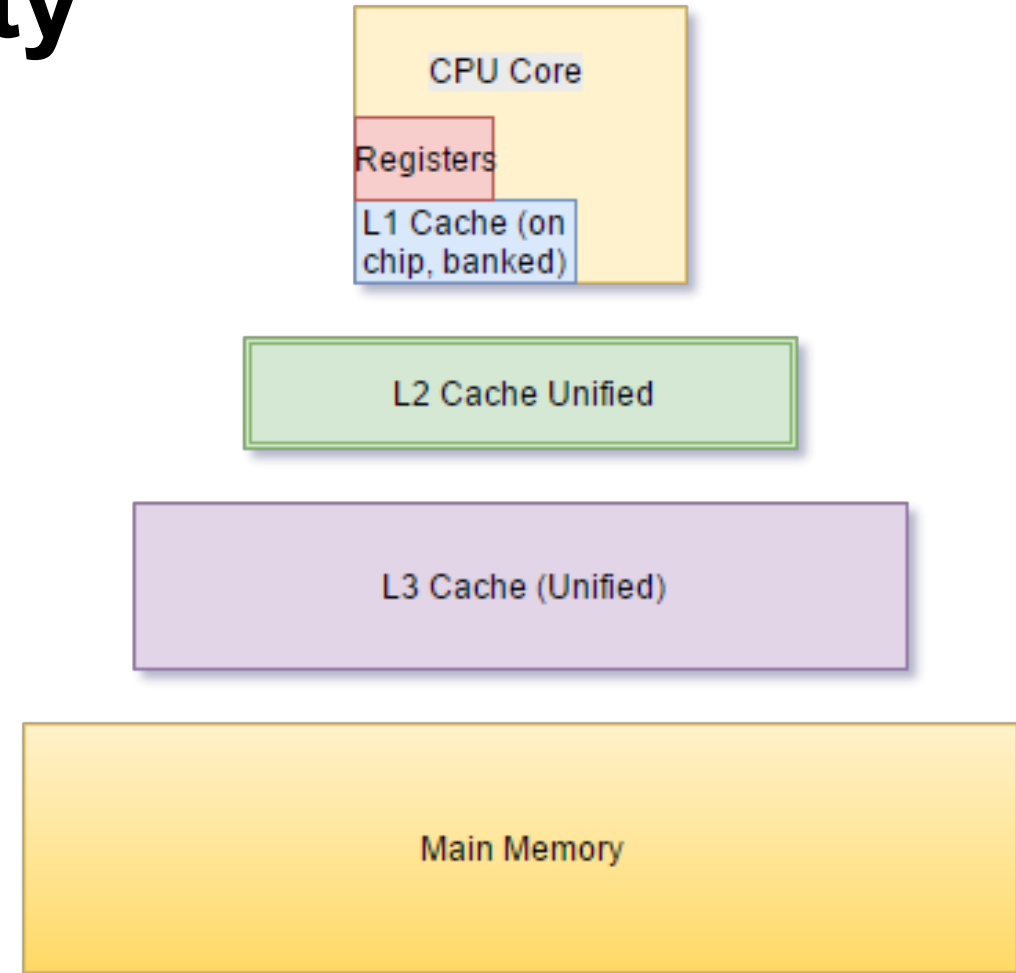
Improving Cache Performance

Reduce miss rate (**Associative cache**)
Reduce miss (time) penalty (??)
Speed hit access time!



Reducing the miss penalty

- Approach: ***Multi-level Cache***
 - Level-1 cache (Primary cache): service the CPU, Small, but fast
 - Level-2 cache: services misses from primary cache, still faster than main memory



Multilevel Cache Example

- Given
 - CPU base CPI = 1, clock rate = 4GHz
 - Miss rate/instruction = 2%
 - Main memory access time = 100ns
- With just a primary cache
 - Miss penalty = $100\text{ns} / 0.25\text{ns} = 400$ cycles
 - Effective CPI = Base CPI + Memory-stall cycles per inst.
 $= 1 + 0.02 \times 400 = 9$

Example (cont.)




- **Now add L-2 cache**
 - Access time = 5ns
 - Global miss rate to main memory = 0.5%
- Primary miss with L-2 hit
 - Penalty = $5\text{ns}/0.25\text{ns} = 20$ cycles
- Total CPI = 1 + primary stall + secondary stall
 - $\text{CPI} = 1 + 0.02 \times 20 + 0.005 \times 400 = 3.4$
- Performance ratio = $9/3.4 = 2.6$
 - Faster by a factor of 2.6 with secondary cache

Multilevel Cache Considerations

- Primary cache
 - Focus on minimal hit time
- L-2 cache
 - Focus on low miss rate to avoid main memory access
 - Hit time has less overall impact
- Results
 - L-1 cache usually smaller than a single cache
 - L-1 block size smaller than L-2 block size

Interactions with Software

- Misses depend on memory access patterns
 - Algorithm behaviour
 - Compiler optimization for memory access



Matrix storage

- Store two-dimensional data in one dimensional (linear) memory



FP Example: Array Multiplication

- MIPS code:

| |
|---|
| li \$t1, 32 # \$t1 = 32 (row size/loop end) |
| li \$s0, 0 # i = 0; initialize 1st for loop |
| L1: li \$s1, 0 # j = 0; restart 2nd for loop |
| L2: li \$s2, 0 # k = 0; restart 3rd for loop |
| sll \$t2, \$s0, 5 # \$t2 = i * 32 (size of row of x) |
| addu \$t2, \$t2, \$s1 # \$t2 = i * size(row) + j |
| sll \$t2, \$t2, 3 # \$t2 = byte offset of [i][j] |
| addu \$t2, \$a0, \$t2 # \$t2 = byte address of x[i][j] |
| l.d \$f4, 0(\$t2) # \$f4 = 8 bytes of x[i][j] |
| L3: sll \$t0, \$s2, 5 # \$t0 = k * 32 (size of row of z) |
| addu \$t0, \$t0, \$s1 # \$t0 = k * size(row) + j |
| sll \$t0, \$t0, 3 # \$t0 = byte offset of [k][j] |
| addu \$t0, \$a2, \$t0 # \$t0 = byte address of z[k][j] |
| l.d \$f16, 0(\$t0) # \$f16 = 8 bytes of z[k][j] |

...

FP Example: Array Multiplication

...

```
sll $t0, $s0, 5      # $t0 = i*32 (size of row of y)
addu $t0, $t0, $s2    # $t0 = i*size(row) + k
sll $t0, $t0, 3       # $t0 = byte offset of [i][k]
addu $t0, $a1, $t0    # $t0 = byte address of y[i][k]
l.d  $f18, 0($t0)     # $f18 = 8 bytes of y[i][k]
```

```
mul.d $f16, $f18, $f16 # $f16 = y[i][k] * z[k][j]
add.d $f4, $f4, $f16   # f4=x[i][j] + y[i][k]*z[k][j]
```

```
addiu $s2, $s2, 1      # $k k + 1
bne  $s2, $t1, L3      # if (k != 32) go to L3
s.d  $f4, 0($t2)       # x[i][j] = $f4
```

```
addiu $s1, $s1, 1      # $j = j + 1
bne  $s1, $t1, L2      # if (j != 32) go to L2
```

```
addiu $s0, $s0, 1      # $i = i + 1
bne  $s0, $t1, L1      # if (i != 32) go to L1
```

Reading

- Sections 5.4