

Operating systems II

CS 2506

Dr. Dan Grigoras
Computer Science Department

Today's topics

- The role of the operating system.
- What is the relationship between the applications/programs and “the computing system”?
- What do we need to learn and why?

Computing systems


- What is a computer now?
- Computer configuration and architecture changed:
 - CPU/processor/core → many cores, homogeneous or heterogeneous
 - Main memory (plain) → memory hierarchy (cache, main memory)
 - Storage (local or remote), disk → array of large disks
 - Bus → multiple buses
 - Peripherals → + Sensors
- ...or a cloud instance provided as a service.
- *Question*: how are these resources allocated to applications?

Allocation of computing resources

- Generally, there is more than one physical unit of a type of resource in a computer configuration. Examples: cores, memory pages, disk sectors, networking cards, cameras.
- Resource allocation to processes is carried out in terms of units and managed by specific OS services during processes life time.
- Many operating systems are using virtualization of resources (cores, memory pages) for better utilization.

Computing resources management

- OS goals:
 - To meet user requirements
 - Least execution time, quick response;
 - Least cost;
 - Fairness;
 - Best possible user experience.
 - Security.
 - To meet system admin requirements
 - Optimal use of resources;
 - Least energy use (green computing);
 - Maximise revenue.
- The system admin can be the user/owner but not necessarily in all cases – see rack-scale and cloud computing.



Computing system layers

- The bottom layer is the hardware; it accepts primitive commands such as “seek the disk arm to track 79, select head 3 and read sector 5”. Software that interacts directly with the hardware is non-portable and contains low level details: data for control and state registers, interrupt priorities, DMA starting address,...
- The OS kernel has several key components:
 - *Device drivers* are the hardware units managers; they hide the low level details.
 - *File sys manager* is the code that organizes the data storage into files and directories, hiding low level details re disk blocks, for example.
 - *Process management* handles processes, allocating them resources and scheduling their execution.
 - *Memory management* is responsible for physical memory and virtual memory management.
 - *Network services* provide host-to-host and process-to-process communication across network.

Access to kernel services

- The kernel can be viewed as a collection of services that user programs may call. They offer functionality and a higher level of abstraction of the computer.
- The repertoire of commands supported by the kernel defines the “virtual machine” which is platform-independent.
- To enter the kernel, the user program makes a *system call by executing a trap instruction of some sort*.
- This instruction switches the processor into a privileged operating mode (*kernel mode*) and jumps into the kernel through a well-defined trap address.
- Parameters passed with the trap instruction tell the kernel what service is requested.
- When the function is completed, the processor flips back to its normal operating mode (*user mode*) and returns control to the user program.

Trap instructions

- There are functions that require specific knowledge of handling resources – control registers, state register, sequence of operations, and a certain degree of protection – resources shared by several users/programs.
- These functions are coded as service routines; they are also known as *system calls*.
- The sequence of steps for a system call is:
 - The system call is invoked by the user program;
 - OS function is performed;
 - Control returns to the user program.
- Trap instructions are used to implement system calls.

Library functions

- User programs have access to libraries and include in their code functions of these libraries – linked in the executables.
- Some library functions use system calls. The function itself parcels up the parameters correctly and then performs the trap.
- The function works as a wrapper for the system call.
- Example:

`printf(“hello world”); ➔ write(1, “hello world”, 11);`

OS classes

- There isn't one OS that fits all computing devices!
- The large range of computing devices requires customized OS.
- General purpose computers run complex OS: Unix, Linux, Windows,...
- Mobile devices, such as smart phones or sensors, run OS like iPhone OS, Android, TinyOS, concerned with power saving and user experience.
- Embedded systems run scaled down versions of OS, real-time, event-driven.

Trends

- Rapid advances in hardware:
 - peripheral devices are integrated onto a die, each with complex, varying programming models, that execute specialized firmware with little OS integration.
 - for energy saving, cores, devices, and memory can be powered on/off at any point during execution.
 - large, non-volatile main memories will likely become prevalent and might have greater capacity than today's disks and disk arrays.
- Applications are changing:
 - many applications use cloud (data center) services, individual users' mobile devices, and sensors and actuators in the Internet of Things.
 - enterprise applications such as file servers, relational databases, and big data analytics now come pre-packaged in a rack-scale software appliance, delivered to customers who have only to plug in and use the system.
 - Clouds are running applications in containers with their own OS.

Suggested reading

- D. Milojevic and T. Roscoe: *Will OSs in 2025 still resemble the Unix-like consensus of today, or will a very different design achieve widespread adoption?*, Computer, Volume: 49, Jan. 2016, pp. 43-51.
- Available in Boole Library:
<http://ieeexplore.ieee.org.ucc.idm.oclc.org/document/7383138>
- My comment: good personal overview of existing operating systems and technology advances that will have a great impact on operating systems architecture and operation. The paper can suggest interesting avenues of further study and research.

Course goals and methodology

- Goals
 - learn the main services of an operating system:
 - processes and threads management, including scheduling and load balancing, IPC;
 - physical and virtual memory management;
 - device drivers;
 - file management system.
 - learn features of different models of operating systems:
 - general purpose;
 - mobile;
 - sensor.
- Methodology
 - attending the lectures.
 - carrying out the lab work.
 - using recommended references to learn more about specific topics.

- Text book:

A. S. Tanenbaum and H. Bos: Modern Operating Systems, Pearson, 4th edition, 2014.

Course philosophy:

Collaborative learning process

Resources: cs4.ucc.ie/moodle

- Grading
 - Continuous assessment: 20% Labs
 - Written exam: 90 min 80%
- Lecture 50 min + 5 min review & questions
- Office: G65, Western Gateway Bldg
- Email: grigoras@cs.ucc.ie, Subject: CS2506
- Office hours: by appointment

Plagiarism

1. Plagiarism is presenting someone else's work as your own. It is a violation of UCC Policy and there are strict and severe penalties.
2. You must read and comply with the UCC Policy on Plagiarism www.ucc.ie/en/exams/procedures-regulations/
3. The Policy applies to *all* work submitted, including software.
4. You can expect that your work will be checked for evidence of plagiarism or collusion.
5. In some circumstances it may be acceptable to reuse a small amount of work by others, but *only* if you provide explicit acknowledgement and justification.
6. If in doubt ask your module lecturer *prior* to submission. Better safe than sorry!