

OVERVIEW OF SQL

- common data sublanguage employed by most relational database management systems (Oracle, Sybase, DB-2, MySQL, SQLServer, Informix, ...)

DDL

CREATE
DROP
ALTER

DML

SELECT
DELETE
UPDATE
INSERT

DCL

GRANT
REVOKE

- the above is the minimum & necessary instruction set for SQL; a particular implementation may support some other commands – and may provide other commands for user convenience
- SQL has undergone several standardization efforts: SQL-1, SQL-2, SQL-86, SQL-3, SQL-92, SQL-99; different DBMS have adhered to different standards, but often with extensions

DATA MANIPULATION LANGUAGE (DML)

```
SELECT [DISTINCT] ColumnList
FROM TableList
[ WHERE Condition ]
[ GROUP BY ColumnList
[HAVING condition ]]
[ ORDER BY ColumnList ]
```

I. BASIC DATA RETRIEVAL (SELECT)

Find the identifier of each project on which some employee has spent 10 or more hours

```
SELECT Pno
FROM Works_On
WHERE Hours >= 10.0;
```

```
SELECT DISTINCT Pno
FROM Works_On
WHERE Hours > 10.0;
```

Find the identifier and name of female employees living in Houston & earning more than 30000.

```
SELECT Ssn, Lname, Fname
FROM Employee
WHERE Sex = 'F'
      AND Salary > 30000
      AND Address LIKE '%Houston%'
ORDER BY Lname, Fname;
```

```
SELECT Ssn, Lname, Fname
FROM Employee
WHERE Sex = 'F'
      AND Salary > 30000
      AND Address LIKE '%Houston%'
ORDER BY Lname DESC, Fname DESC;
```

Note:

The primary issue here is how to specify logical/Boolean conditions [in the WHERE clause] that identify what data is required by a user:

ColumnName operator ColumnName

ComumnName operator Value

where the Boolean operator falls into one of the categories below:

Comparison Operator:

=	<> (!=)	>	<	>=	<=
(NOT =	NOT <>	NOT <	NOT >	NOT >=	NOT <=)

Set Inclusion Operator:

IN (NOT IN)

Subrange Operator:

BETWEEN (NOT BETWEEN)

Substring Operator:

LIKE (NOT LIKE)

Null Test Operator:

IS NULL (IS NOT NULL)

Boolean conditions can be negated by means of the NOT operator and can be combined by means of the AND and OR operators. As with general programming languages, NOT takes precedence over AND which takes precedence over OR. This precedence of operators can be over-ridden by means of parentheses.

Find the name of employees earning between \$20,000 and \$30,000

```
SELECT Lname, Fname
FROM EMPLOYEE
WHERE Salary >= 20000
      AND Salary <= 30000;
```

```
SELECT Lname, Fname
FROM EMPLOYEE
WHERE Salary BETWEEN 20000 AND 30000;
```

Find the name of male employees aged between 21 and 25

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE Sex = 'M'
      AND DOB >= '1981-01-20'
      AND DOB <= '1986-01-19';
```

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE Sex = 'M'
      AND DOB BETWEEN '1981-01-20' AND '1986-01-19';
```

Find the identifier and name of employees earning less than \$25,000 who work in departments 4 or 5

```
SELECT Ssn, Lname, Fname
FROM EMPLOYEE
WHERE Salary < 25000
      AND ( Dno = 4
            OR Dno = 5 )
```

```
SELECT Ssn, Lname, Fname
FROM EMPLOYEE
WHERE Salary < 25000
      AND Dno IN (4, 5)
```

Find the identifier and name of employees earning less than \$25,000 who do not work in department nos. 4 or 5

```
SELECT Ssn, Lname, Fname
FROM Employee
WHERE Salary < 25000
      AND Dno != 4
      AND Dno != 5
```

```
SELECT Ssn, Lname, Fname
FROM Employee
WHERE Salary < 25000
      AND Dno NOT IN (4, 5)
```

```
SELECT Ssn, Lname, Fname
FROM Employee
WHERE Salary < 25000
      AND NOT ( Dno = 4
                OR Dno = 5 )
```

Find the name and address of employees of department 4 with no specified supervisor

```
SELECT Lname, Fname, Address
FROM EMPLOYEE
WHERE Dno = 4
      AND SuperSsn IS NULL
ORDER BY Lname DESC, Fname DESC
```

Note:

Most versions of SQL support *string expressions and functions* for usefulness.

e.g. in Oracle:

```
SELECT Ssn AS "Staff Id", Lname + ', ' + Fname AS "Staff Name"
FROM Employee
WHERE Sex = 'F'
      AND Dno = 5;
```

or, in MySQL:

```
SELECT Ssn AS "Staff Id", CONCAT(Lname, ', ', Fname) AS "Staff Name"
FROM Employee
WHERE Sex = 'F'
      AND Dno = 5;
```

Some useful MySQL string functions are:

CONCAT (str1, str2, ...):	concatenate strings
LEFT (str, len):	returns leftmost <i>len</i> characters from string <i>str</i>
LENGTH (str):	returns string length
LOWER (str):	returns lower-case version of <i>str</i>
REVERSE (str):	reverse the string <i>str</i>
RTRIM (str):	removes trailing blanks
SOUNDEX (str):	returns phonetic sequence of <i>str</i>

UPPER (str): returns upper-case version of *str*

Find the identifier and name of every employee named Smith

SELECT Ssn, UPPER(Lname), Fname	SELECT Ssn, UPPER(Lname), Fname
FROM Employee	FROM Employee
WHERE Lname = 'Smith';	WHERE UPPER(Lname) = 'SMITH';

Find the identifier and name of every employee whose name sounds like Smith [Smith, Smyth, Smythe, Smit, ...]

```
SELECT Ssn AS "Employee Id", CONCAT(UPPER(Lname), ' ', Fname) AS "Employee Name"
FROM Employee
WHERE SOUNDEX(Lname) = SOUNDEX('Smith');
```

Note:

Most versions of SQL also support *numeric expressions and functions* for usefulness.

Arithmetic Operators: + - * /

e.g.

```
SELECT Salary/12
FROM Employee
WHERE Ssn = '123456789'
```

e.g.

```
SELECT Ssn, Salary * 1.06 AS "Employee Cost"
FROM Employee
```

e.g.

```
SELECT Lname + ' ', ' ' + Fname AS "Staff Name", Salary / 12 AS "Monthly Salary"
FROM Employee
WHERE Salary / 12 > 2000
```

```
SELECT Lname + ' ', ' ' + Fname AS "Staff Name", FLOOR(Salary / 12) AS "Monthly Salary"
FROM Employee
WHERE Salary / 12 > 2000
```

Some useful MySQL numeric functions are:

ABS(x):	absolute value of <i>x</i>
CEILING(x), FLOOR(x):	smallest/largest integer closest to value <i>x</i>
COS(x), SIN(x), TAN(x):	sinusoidal functions of <i>x</i>
EXP (x):	exponential of <i>x</i>
LOG (x):	logarithm of value <i>x</i>

Find the name of male employees aged between 21 and 25 (again)

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE Sex = 'M'
```

```
AND DOB BETWEEN DATE_SUB(CURDATE(), INTERVAL 25 YEAR)
AND DATE_SUB(CURDATE(), INTERVAL 21 YEAR);
```

II. STATISTICAL DATA RETRIEVAL (FUNCTIONS, GROUP BY, HAVING)

Find the highest salary of all employees

```
SELECT MAX(Salary)
FROM EMPLOYEE
```

Find the average salary for female employees of department no. 5

```
SELECT AVG(Salary)
FROM EMPLOYEE
WHERE Dno = 5
AND Sex = 'F'
```

Find the total number of employees that have worked on project no. 20, together with the total number of hours expended on it

```
SELECT COUNT(*) AS "Employee Count", SUM(Hours) AS "Hours Expended"
FROM WORKS_ON
WHERE Pno = 20
```

Find the total number of employees who have female dependents

```
SELECT COUNT (DISTINCT Essn)
FROM DEPENDENT
WHERE Sex = 'F'
```

Note:

- the group functions differ from the string and numeric functions seen earlier: the latter apply to a specific value(s) in a specific row of a table; the group functions apply to the entire column(s) of the result table of a query
- either '*' or 'DISTINCT Colname' are the standard arguments to the COUNT function
- the group functions are for computation of basic statistical values; it does not make sense – nor is it legal – to mix normal columns and group values in the same SELECT clause [with one exceptional case, as outlined below]
- the group functions are for computation of basic statistical values; it does not make sense – nor is it legal – to mix normal columns and group values in the same SELECT clause [with one exceptional case, as outlined below]

The useful MySQL group functions are:

AVG(<i>expr</i>):	average value of <i>expr</i>
COUNT(*):	number of rows in result table
COUNT(DISTINCT <i>expr</i>):	number of distinct rows containing <i>expr</i>
MAX(<i>expr</i>):	maximum value of <i>expr</i>
MIN(<i>expr</i>)	minimum value of <i>expr</i>
SUM (<i>expr</i>):	summation of all values of <i>expr</i>

The *expr* value of AVG & SUM must be numeric

Find the highest salary in each department

```
SELECT Dno, MAX(Salary)
FROM EMPLOYEE
GROUP BY Dno
```

Find the average monthly salary for female employees in each department

```
SELECT Dno, AVG(Salary / 12)
FROM EMPLOYEE
WHERE Sex = 'F'
GROUP BY Dno
```

Note:

- this last example is the exceptional case, where normal columns and group functions can be used together in SELECT clauses: the column concerned is the one on which grouping occurs in the GROUP BY clause
- the effect of the GROUP BY clause is to:
 - arrange the table (after restriction as defined by the WHERE clause) into similarity groups of rows that match on the specified column(s);
 - treat each group of rows as though it was a separate table for group function calculation]

Find the highest salary in each department with more than five employees

```
SELECT Dno, MAX(Salary)
FROM EMPLOYEE
GROUP BY Dno
HAVING COUNT(*) >= 5
```

Find the identity of each department, and its average salary, whose average salary is greater than \$35,000

```
SELECT Dno, AVG(Salary)
FROM EMPLOYEE
```

GROUP BY Dno

HAVING AVG(Salary) > 35000

III. MULTIPLE TABLE RETRIEVAL USING SUBQUERIES

Find the name of employees working for the Research Department

```
SELECT CONCAT (UPPER(Lname), ' ', Fname) AS Name
FROM Employee
WHERE Dno IN
    ( SELECT Dnumber
      FROM DEPARTMENT
      WHERE Dname = 'Research' )
```

Find the name of male employees working in Houston

```
SELECT CONCAT (Lname, ' ', Fname) AS Name
FROM EMPLOYEE
WHERE Sex = 'M'
    AND Dno IN
        ( SELECT Dnumber
          FROM DEPT_LOCATIONS
          WHERE Dlocation = 'Houston' )
ORDER BY Name
```

Find the identifier and address of female employee working on either 'ProductX' or 'ProductY' projects

```
SELECT Ssn , Address
FROM EMPLOYEE
WHERE Sex = 'F'
    AND Ssn IN
        ( SELECT Essn
          FROM WORKS_FOR
          WHERE Pno IN
              ( SELECT Pnumber
                FROM PROJECT
                WHERE Pname IN ('ProductX', 'ProductY' ) )
```

Find the name(s) of the highest earner / highest female earner

```
SELECT Lname, Fname
FROM EMPLOYEE
WHERE Salary =
  ( SELECT MAX (Salary)
    FROM EMPLOYEE )
```

```
SELECT Lname, Fname
FROM EMPLOYEE
WHERE Salary >=ALL
  ( SELECT Salary
    FROM EMPLOYEE )
```

```
SELECT Lname, Fname
FROM EMPLOYEE
WHERE Sex = 'F'
AND Salary =
  ( SELECT MAX (Salary)
    FROM EMPLOYEE
  WHERE Sex = 'F' )
```

```
SELECT Lname, Fname
FROM EMPLOYEE
WHERE Sex = 'F'
AND Salary >=ALL
  ( SELECT Salary
    FROM EMPLOYEE
  WHERE Sex = 'F' )
```

Interblock Connectives

An interblock connective is an operator in the WHERE clause that expresses a condition between a column value and the result returned by a subquery. In the examples above, we used IN, = and >=ALL as interblock connectives. A full list is given below:

Comparison Operators:

= != > < >= <=

Set Comparison Operators:

=ANY <>ANY >ANY <ANY >=ANY <=ANY
=ALL <>ALL >ALL <ALL >=ALL <=ALL

Set Inclusion Operator:

IN (NOT IN)

Existential Quantifier:

EXISTS (NOT EXISTS)

Notes:

1. SQL is said to be block structured in that a query can be organised into blocks (or levels, or subqueries). Each level can be an arbitrary query in its own right and can use any of the features of the SQL language, with three exceptions: the SELECT DISTINCT option, the use of the ORDER BY clause and the renaming of retrieved columns are only permitted at the outer level. Note that all references to columns are local to the current block's FROM clause.
2. All subqueries, except for synchronised subqueries to be seen later, evaluate bottom-up, i.e., each individual level executes to completion and returns a result to the next level up, where it can be used in testing a condition.

3. The comparison operators are valid only if the subquery returns a single row. Otherwise, an error will occur. The set comparison operators should be used when a set of values are expected back from the subquery.
4. When using the set comparison operators, 'ANY' corresponds to common usage of the word some; similarly, 'ALL' would generally be expressed in English as every. Thus, in the last example above, we are searching for the EMPLOYEE member whose salary is greater than or equal to every salary value in existence; obviously, therefore, it must be the largest such value.
5. =Any and IN are equivalent. Similarly, <>ALL is equivalent to NOT IN. =ALL is redundantly false and should not be used; the same goes for <>ANY, which is redundantly true. Thus, the two queries shown below are of no real use.
6. The existential quantifier is used only with a particular type of subquery – synchronised subqueries – which we shall see later.

Redundant Queries:

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE Ssn =ALL
  ( SELECT Essn
    FROM WORKS_ON
    WHERE Hours >= 10.0 )
```

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE Ssn <>ANY
  ( SELECT Essn
    FROM WORKS_ON
    WHERE Hours >= 10.0 )
```

Find the name & address of the lowest-paid employee living in Houston

```
SELECT Lname, Fname, Address
FROM EMPLOYEE
WHERE Address LIKE '%Houston%'
  AND Salary =
    ( SELECT MIN (Salary)
      FROM EMPLOYEE
      WHERE Address LIKE '%Houston%' )
```

or

```
SELECT Lname, Fname, Address
FROM EMPLOYEE
WHERE Address LIKE '%Houston%'
  AND Salary <= ALL
    ( SELECT Salary
      FROM EMPLOYEE
      WHERE Address LIKE '%Houston%' )
```

Find the identifier and name of projects in Bellaire that were never worked on

```
SELECT Pno, Pname
FROM PROJECT
WHERE Plocation = 'Bellaire'
  AND Pnumber NOT IN
    ( SELECT Pno
      FROM WORKS_ON )
```

Find the identifier, name and location of the project(s) which has had the most people working on it

```
SELECT Pnumber, Pname, Plocation
FROM PROJECT
WHERE Pnumber IN
    ( SELECT Pno
      FROM WORKS_ON
      GROUP BY Pno
      HAVING COUNT (*) = MAX (COUNT (*) ) )
```

Note:

This last query, while adhering to the SQL standard, will generate an error in MySQL. This is because MySQL does not support embedded function calls – the calculation of a function of a function.

Can you figure out how to rewrite the query so that it does work in MySQL? Hint: You will need one additional level subquery.

Find the name of employees who have the same department and salary as 123456789

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE Ssn <> '123456789'
    AND (Dno, Salary) =
        ( SELECT Dno, Salary
          FROM EMPLOYEE
          WHERE Sno = '123456789' )
```

Note:

This query will also generate an error in MySQL, because that DBMS only permits a subquery to return a single column. It could be rewritten using independent subqueries as shown below:

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE Ssn <> '123456789'
    AND Dno =
        ( SELECT Dno
          FROM EMPLOYEE
          WHERE Ssn = '123456789' )
    AND Salary =
        ( SELECT Salary
          FROM EMPLOYEE
          WHERE Ssn = '123456789' )
```

However, this technique has its limits. The following query – which purports to find the name of employees who have the same department and salary as John Smith – would produce incorrect results if more than one John Smith existed. Can you see why?

```
SELECT Lname + ', ' + Fname
FROM EMPLOYEE
WHERE ( Fname <> 'John' OR Lname <> 'Smith' )
    AND Dno IN
        ( SELECT Dno
          FROM EMPLOYEE
          WHERE Fname = 'John'
            AND Lname = 'Smith' )
```

```

AND Salary IN
( SELECT Salary
  FROM EMPLOYEE
  WHERE Fname = 'John'
    AND Lname = 'Smith' )

```

IV. MULTIPLE TABLE RETRIEVAL USING JOINS

Find the name of employees working for the Research Department

```

SELECT CONCAT (UPPER(Lname), ' ', Fname) AS Name
FROM EMPLOYEE, DEPARTMENT
WHERE Dno = Dnumber
    AND Dname = 'Research'
ORDER BY Name;

```

In MySQL, the following is also allowed:

```

...
FROM EMPLOYEE JOIN DEPARTMENT
...

```

Find the name of male employees working in Houston

```

SELECT CONCAT (Lname, ' ', Fname) AS Name
FROM EMPLOYEE, DEPT_LOCATIONS
WHERE Sex = 'M'
    AND Dno = Dnumber
    AND Dlocation = 'Houston'
ORDER BY Name;

```

Find the identifier and address of female employee working on either 'ProductX' or 'ProductY' projects

```

SELECT Ssn , Address
FROM EMPLOYEE, WORKS_ON, PROJECT
WHERE Ssn = Essn
    AND Pno = Pnumber
    AND Pname IN ('ProductX', 'ProductY' )
    AND Sex = 'F';

```

or

```

SELECT Ssn , Address
FROM EMPLOYEE
WHERE Ssn IN
( SELECT Essn
  FROM WORKS_ON , PROJECT
  WHERE Pno = P.number
    AND Pname IN ('ProductX', 'ProductY' ))
AND Sex = F;

```

Find the name(s) & departments of the highest earner / highest female earner

```
SELECT Lname, Fname, Dname
FROM EMPLOYEE, DEPARTMENT
WHERE Dno= Dnumber
      AND Salary =
      ( SELECT MAX (Salary)
        FROM EMPLOYEE )
```

```
SELECT Lname, Fname, Dname
FROM EMPLOYEE, DEPARTMENT
WHERE Dno = Dnumber
      AND Sex = 'F'
      AND Salary =
      ( SELECT MAX (Salary)
        FROM STAFF
        WHERE Sex = 'F' )
```

Find the name of employees working in the same department as Joyce English

```
SELECT RE.Fname, RE.Lname
FROM EMPLOYEE LE, EMPLOYEE RE
WHERE LE.Dno = RE.Dno
      AND LE.Ssn != RE.Ssn
      AND LE.Fname = 'Joyce'
      AND LE.Lname = 'English'
```

Find the name of employees working in the same department as Joyce English, but earning more than her

```
SELECT RE.Fname, RE.Lname
FROM EMPLOYEE LE, EMPLOYEE RE
WHERE LE.Dno = RE.Dno
      AND LE.Ssn != RE.Ssn
      AND LE.Fname = 'Joyce'
      AND LE.Lname = 'English'
      AND LE.Salary < RE.Salary
```

Note:

The second condition above (LE.Ssn != RE.Ssn) is actually redundant in light of the final condition (LE.Salary < RE.Salary)

Find the name of every Texax-based employee, together with the name & relationship of any dependents he/she might have

```
SELECT Fname, Lname, Dependent_Name, Relationship
FROM EMPLOYEE LEFT JOIN DEPENDENT
WHERE Ssn = Essn
      AND Address LIKE '%TX'
```

- this requests that a *left outer join* be evaluated [any normally unmatched row from the left table will automatically match a null row appended to the right]
- a right outer join also exists in MySQL

- the syntax user here is MySQL-specific; other implementations of the SQL language adopt different syntax

V. MULTIPLE TABLE RETRIEVAL USING SYNCHRONIZED SUBQUERIES

also known as Correlated Subqueries

also known as Subqueries with Interblock Reference

Find the name of employees working in the Research department

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE 'Research' IN
  ( SELECT Dname
    FROM DEPARTMENT
    WHERE Dnumber = Dno )
```

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE 'Research' =
  ( SELECT Dname
    FROM DEPARTMENT
    WHERE Dnumber = Dno )
```

Note:

The two above queries, while adhering to the SQL language standard, will not be accepted by MySQL. To legitimize them, we must explicitly alias the outer table, and then use that alias within the subquery:

```
SELECT Fname, Lname
FROM EMPLOYEE E
WHERE 'Research' IN
  ( SELECT Dname
    FROM DEPARTMENT
    WHERE Dnumber = E.Dno )
```

```
SELECT Fname, Lname
FROM EMPLOYEE E
WHERE 'Research' =
  ( SELECT Dname
    FROM DEPARTMENT
    WHERE Dnumber = E.Dno )
```

Find the name of male employees working in Houston

```
SELECT Fname, Lname
FROM EMPLOYEE
WHERE Sex = 'M'
  AND 'Houston' IN
    ( SELECT Dlocation
      FROM DEPT_LOCATIONS
      WHERE Dnumber = Dno )
ORDER BY Name
```

Find the identifier & name of departments hiring employees resident in Houston

```
SELECT Dnumber, Dname
FROM DEPARTMENT D
WHERE EXISTS
  ( SELECT *
    FROM EMPLOYEE
    WHERE Address LIKE '%Houston%'
      AND Dno = D.Dnumber )
```

Note:

The EXISTS condition returns a *true* value if the subquery returns any data; otherwise it returns *false*. Correspondingly, NOT EXISTS returns *true* when the subquery returns nothing.

Find the identifier & name of departments hiring no employees resident in Houston

```
SELECT Dnumber, Dname
FROM DEPARTMENT D
WHERE NOT EXISTS
  ( SELECT *
    FROM EMPLOYEE
    WHERE Address LIKE '%Houston%'
      AND Dno = D.Dnumber )
```

Find the name of all employees working on project no. 20

```
SELECT Fname, Lname
FROM EMPLOYEE E
WHERE EXISTS
  ( SELECT *
    FROM WORKS_ON
    WHERE Essn = E.Ssn
      AND Pno = 20 )
```

Find the name of staff handling working on no projects

```
SELECT Fname, Lname
FROM EMPLOYEE E
WHERE NOT EXISTS
  ( SELECT *
    FROM WORKS_ON
    WHERE Essn = E.Ssn )
```

VI. DATA INSERTION, REMOVAL & MODIFICATION

Note:

Unlike retrieval (SELECT statement), the other DML statements (INSERT, DELETE & UPDATE) can potentially leave the database in an inconsistent state. For example, inserting a row into the EMPLOYEE table with a Dno value of 50 is in conflict with the fact that no DEPARTMENT row with Dnumber 50 exists. Similar problems can arise due to deletion or modification of values that implicitly link tables together. Such inconsistencies are termed *referential integrity violations*.

Database administrators might address this issue in one of three ways:

- (i) Do Nothing [or almost nothing].
Database modification might be confined to a small number of highly trained individuals, who are then trusted to do things properly.
Advantage: Easy to do.
Disadvantage: Corruption of the database is inevitable.
- (ii) Specify rules that the DBMS might apply to modification commands.
DBMS supports *constraint rules* that can prevent data corruption.
Advantage: Automated consistency checks are possible.

Disadvantage: Will slow down evaluation of SQL data modification commands; handling of problematic commands might not be straightforward; reasons for command rejection might not be obvious to users.

- (iii) Perform data modification under program control.
Do not directly use SQL commands for database modification; instead, embed the commands within a programming language (e.g. PHP) that also carries out sanity checks on what the commands are attempting to do.
Advantage: Can apply sophisticated command validation, standards application, etc.
Disadvantage: Requires the writing of programs to handle multiple common data modification processes.

Record details of a new employee.

```
INSERT INTO EMPLOYEE
VALUES ( 'Joan', 'J', 'McGregor', '234765980', '1978-01-01', '81 Beech, Austin, TX', 'F', 75000,
NULL, 4);
```

Update details of an employee.

```
UPDATE EMPLOYEE
SET Address = '18 Maple, Austin, TX',
    Salary = 50000
WHERE Ssn = '123456789';
```

Increment the salary of employees of the Research department by 10%

<pre>UPDATE EMPLOYEE SET Salary = Salary * 1.1 WHERE Dno IN (SELECT Dnumber FROM DEPARTMENT WHERE Dname = 'Research');</pre>	<pre>UPDATE EMPLOYEE SET Salary = Salary * 1.1 FROM EMPLOYEE, DEPARTMENT WHERE Dno = Dnumber AND Dname = 'Research';</pre>
--	--

Remove information on Employee Jennifer Wallace.

```
DELETE
FROM EMPLOYEE
WHERE FName = 'Jennifer'
    AND LName = 'Wallace';
```

Remove information on all dependents of Franklin Wong.

<pre>DELETE FROM DEPENDENT WHERE Essn IN (SELECT Ssn FROM EMPLOYEE WHERE FName = 'Franklin' AND LName = 'Wong');</pre>	<pre>DELETE DEPENDENT FROM EMPLOYEE, DEPENDENT WHERE Ssn = Essn AND FName = 'Franklin' AND LName = 'Wong';</pre>
--	--

Note:

We can combine INSERT with SELECT to populate one table with data extracted from another. This is a useful feature for data extraction/restructuring.

e.g. assume we have a newly-created (empty) table:

HOUSTON_STAFF (StaffNo, Name, Dept, Salary)

We could populate this table using existing contents of the EMPLOYEE table:

```
INSERT INTO HOUSTON_STAFF
SELECT Ssn, CONCAT (LName, ' ', FName, ' ', Minit), Dno, Salary
FROM EMPLOYEE
WHERE Address LIKE '%Houston%';
```


DATA DEFINITION LANGUAGE (DDL)

CREATE / DROP DATABASE

```
CREATE DATABASE companydb;
```

```
CREATE DATABASE IF NOT EXISTS mydatabase;
```

```
CREATE DATABASE Dreamhome
```

```
ON D:/dbm2                      { Oracle }
```

```
INITIAL=4, EXTENT=2
```

```
LOG ON D:/dblogs;
```

```
DROP DATABASE companydb;
```

CREATE / DROP / ALTER TABLE

```
CREATE TABLE PROJECT
```

```
(  PNAME      CHAR(15),  
   PNUMBER    INT,  
   PLOCATION   VARCHAR(15),  
   DNUM       INT );
```

```
CREATE TABLE PROJECT
```

```
(  PNAME      VARCHAR(15) NOT NULL,  
   PNUMBER    INT          NOT NULL,  
   PLOCATION   VARCHAR(15),  
   DNUM       INT          NOT NULL,  
   PRIMARY KEY (PNUMBER),  
   UNIQUE (PNAME),  
   FOREIGN KEY (DNUM) REFERENCES DEPARTMENT(DNUMBER) );
```

Note:

The last example shows how a table definition requires a *table name*, *attribute names* and *attribute data types*; it can also contain *constraints* of various types, such as *null constraints*, *key/unique constraints* and *foreign/referential constraints*.

Several variations in syntax are possible, as outlined below:

```
PLOCATION    VARCHAR(15) DEFAULT 'Houston',

CONSTRAINT UNIQUE (PNAME),

CONSTRAINT OneName UNIQUE (PNAME),    { constraint name }
```

```
CREATE TABLE PROJECT
( PNAME          VARCHAR(15) NOT NULL,
  PNUMBER        INT          NOT NULL,
  PLOCATION       VARCHAR(15),
  DNUM INT              NOT NULL,
  PRIMARY KEY (PNUMBER),
  CONSTRAINT OneName UNIQUE (PNAME),
  CONSTRAINT Controller FOREIGN KEY (DNUM) REFERENCES
    DEPARTMENT(DNUMBER) ON DELETE CASCADE );
```

CREATE TABLE EMPLOYEE

(FNAME	VARCHAR(15)	NOT NULL ,
MINIT	CHAR ,	
LNAME	VARCHAR(15)	NOT NULL ,
SSN	CHAR(9)	NOT NULL ,
BDATE	DATE	
ADDRESS	VARCHAR(30) ,	
SEX	CHAR ,	
SALARY	DECIMAL(10,2) ,	
SUPERSSN	CHAR(9) ,	
DNO	INT	NOT NULL ,

PRIMARY KEY (SSN) ,**FOREIGN KEY** (SUPERSSN) **REFERENCES** EMPLOYEE(SSN) ,**FOREIGN KEY** (DNO) **REFERENCES** DEPARTMENT(DNUMBER)) ;**CREATE TABLE DEPARTMENT**

(DNAME	VARCHAR(15)	NOT NULL ,
DNUMBER	INT	NOT NULL ,
MGRSSN	CHAR(9)	NOT NULL ,
MGRSTARTDATE	DATE ,	

PRIMARY KEY (DNUMBER) ,**UNIQUE** (DNAME) ,**FOREIGN KEY** (MGRSSN) **REFERENCES** EMPLOYEE(SSN)) ;**CREATE TABLE DEPT_LOCATIONS**

(DNUMBER	INT	NOT NULL ,
DLOCATION	VARCHAR(15)	NOT NULL ,

PRIMARY KEY (DNUMBER, DLOCATION) ,**FOREIGN KEY** (DNUMBER) **REFERENCES** DEPARTMENT(DNUMBER)) ;

```

CREATE TABLE PROJECT
  ( PNAME          VARCHAR(15)      NOT NULL ,
    PNUMBER        INT              NOT NULL ,
    PLOCATION       VARCHAR(15) ,
    DNUM           INT              NOT NULL ,
    PRIMARY KEY (PNUMBER) ,
    UNIQUE (PNAME) ,
    FOREIGN KEY (DNUM) REFERENCES DEPARTMENT(DNUMBER) ) ;

CREATE TABLE WORKS_ON
  ( ESSN           CHAR(9)          NOT NULL ,
    PNO            INT              NOT NULL ,
    HOURS          DECIMAL(3,1)     NOT NULL ,
    PRIMARY KEY (ESSN, PNO) ,
    FOREIGN KEY (ESSN) REFERENCES EMPLOYEE(SSN) ,
    FOREIGN KEY (PNO) REFERENCES PROJECT(PNUMBER) ) ;

CREATE TABLE DEPENDENT
  ( ESSN           CHAR(9)          NOT NULL ,
    DEPENDENT_NAME VARCHAR(15)     NOT NULL ,
    SEX            CHAR ,
    BDATE          DATE ,
    RELATIONSHIP   VARCHAR(8) ,
    PRIMARY KEY (ESSN, DEPENDENT_NAME) ,
    FOREIGN KEY (ESSN) REFERENCES EMPLOYEE(SSN) ) ;

```

Note:

- the effect of the CREATE TABLE command is to store a descriptor of the table in the *database directory* [system catalog]; the effect of ALTER TABLE is to modify that descriptor; the effect of DROP TABLE is to remove the descriptor
- primary/unique columns must have NOT NULL constraints
- MySQL: constraint names are ignored (cannot be dropped in ALTER TABLE statements)
- MySQL: a single variable-length column, explicit or implicit, causes all columns to be silently converted to variable-length data types

Data Types in MySQL

Numeric	String	Date/Time
TINYINT	CHAR	DATE ('CCYY-MM-DD')
SMALLINT	VARCHAR	TIME ('hh:mm:ss')
MEDIUMINT	TINYBLOB	DATETIME
INT	BLOB	TIMESTAMP
BIGINT	MEDIUMBLOB	YEAR ('CCYY')
FLOAT	LOBLOB	
DOUBLE	TINYTEXT	
DECIMAL	TEXT	
	MEDIUMTEXT	
	LONGTEXT	
	ENUM	
	SET	

Note:

- variants of the same basic type (e.g. INT family) differ in their storage requirements – and upper/lower limits (e.g. 1, 2, 3, 4, 8 bytes, with values up to 127, 32767, 8388607, 2147483647, ...); [actually, these values are for signed numbers, unsigned (positive) numbers can double in size]
- these types may be parameterized for storage and/or display purposes:
 - INT (3) - displays in column width 3
 - DECIMAL (7,2) - displays in “7.2” format
 - CHAR (15) - storage 15 characters, padded
 - VARCHAR (15) - max storage 15 characters, unpadded
- the TEXT & BLOB families are implicitly of variable length
- observations:
 - variable-length types are designed for saving of storage space (!)
 - TEXT/BLOB types are relatively uncommon, as they don't fit easily into a table structure; they are rarely used as conditions in queries; when retrieved within queries – not necessarily as part of the condition – they can prove inefficient; to overcome this, these types are often extracted into a separate table

```
ALTER TABLE PROJECT
ADD START_DATE DATE;
```

```
ALTER TABLE EMPLOYEE
ADD POSITION CHAR(12) NOT NULL;      { legal? }
```

```
ALTER TABLE PROJECT
```

DROP PLOCATION;

ALTER TABLE EMPLOYEE
MODIFY ADDRESS CHAR(50);

DROP TABLE DEPT_LOCATIONS;

CREATE / DROP INDEX

CREATE INDEX EMPX ON EMPLOYEE(SSN);

CREATE INDEX WEP ON WORKS_ON(ESSN, PNO);

CREATE UNIQUE INDEX DEPTX ON DEPARTMENT(DNUMBER);

DROP INDEX EMPX;

Note:

- an index is a secondary storage item that should speed up queries on the indexed column(s); it is created if certain query patterns are expected; it may, however, slow down inserts, deletes & updates on the table; note that database operations can never subsequently request that an index be used during evaluation – it is the task of the *query optimizer* to use or ignore the presence of an index in generating an evaluation plan
- indexes can also be created within the CREATE TABLE or ALTER TABLE statements; they can be dropped within the ALTER TABLE statements
- MySQL: automatically converts CREATE INDEX & DROP INDEX statements into equivalent ALTER TABLE statements

DATA CONTROL LANGUAGE (DCL)

User Account Management

```
GRANT CONNECT  
TO Black  
IDENTIFIED BY BlkPswXX;
```

```
GRANT RESOURCE  
TO white  
IDENTIFIED BY SecretPsw;
```

```
GRANT DBA  
TO SysMgr  
IDENTIFIED BY TopSecrPsw;
```

```
GRANT RESOURCE  
TO Black;
```

```
REVOKE RESOURCE  
FROM Black;
```

Note:

- Creating a new user account requires a username, a password and a privilege level; adding a privilege level to an existing user is also possible
- Privilege levels determine what that user can subsequently do – specifically what SQL commands he/she can legally execute:
 - CONNECT: login to DBMS only
 - RESOURCE: CONNECT + execute DDL statements
 - DBA: RESOURCE + GRANT/REVOKE privilege

Table Access Management

```
GRANT SELECT  
ON companydb.employee  
TO white;
```

```
GRANT SELECT  
ON companydb.*  
TO Black  
WITH GRANT OPTION;
```

```
GRANT SELECT, INSERT
ON companydb.department
TO gray
WITH GRANT OPTION;
```

```
GRANT SELECT, UPDATE (Salary, SuperSSN)
ON companydb.employee
TO Ruby, Green, Black;
```

```
GRANT ALL
ON companydb.*
TO violet
IDENTIFIED BY passwdX
```

```
GRANT ALL
ON companydb.*
TO silver
WITH GRANT OPTION;
```

```
REVOKE SELECT
ON companydb.department
FROM brown;
```

```
REVOKE ALL
ON companydb.*
FROM violet;
```

Note:

- the GRANT & REVOKE commands gives/takes permissions on tables/databases for named database users
- the important permissions are:
 - o ALTER [table], INDEX (create & drop), SELECT, INSERT, DELETE, UPDATE, ALL
- the ON clause can specify:
 - o database.tablename - a specific table
 - o database.* - all tables in a specific database
 - o *.* - all databases, all tables (global)
- note that SQL (including MySQL implementation) is achieved by limiting the commands available to a user – if a user cannot execute a command, then he/she is limited in their access to the data/database

- note also that the MySQL system comes configured for one user (the SuperUser, or DBA) who has all permissions