

CS2505: Transport Layer

Copyright Notice: The CS2505 lecture notes are adapted from material provided by J.F. Kurose and K.W. Ross and includes material by C.J. Sreenan, L.L. Peterson, B.S. Davie and others. This material is copyrighted and so these lecture notes must not be copied or distributed without permission.

University College Cork CS2505: Transport Layer 3-1

Outline

Our goals:

- understand principles behind transport layer services:
 - ❖ Multiplexing & demultiplexing
 - ❖ reliable data transfer
 - ❖ flow control
 - ❖ congestion control
- learn about transport layer protocols in the Internet:
 - ❖ UDP: connectionless transport
 - ❖ TCP: connection-oriented transport
 - ❖ TCP congestion control

University College Cork CS2505: Transport Layer

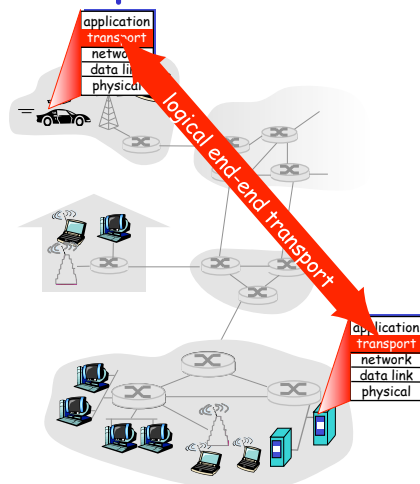
3-2

Outline

- ❑ 3.1 Transport-layer services
- ❑ 3.2 Multiplexing and demultiplexing
- ❑ 3.3 Connectionless transport: UDP
- ❑ 3.4 Principles of reliable data transfer
- ❑ 3.5 Connection-oriented transport: TCP
- ❑ 3.6 TCP congestion control

Transport services and protocols

- ❑ provide *logical communication* between app processes running on different hosts
- ❑ transport protocols run in end systems
 - ❖ send side: breaks app messages into **segments**, passes to network layer
 - ❖ rcv side: reassembles segments into messages, passes to app layer
- ❑ more than one transport protocol available to apps
 - ❖ Internet: TCP and UDP



Transport vs. network layer

- ❑ *network layer*: logical communication between hosts
- ❑ *transport layer*: logical communication between processes
 - ❖ relies on, enhances, network layer services

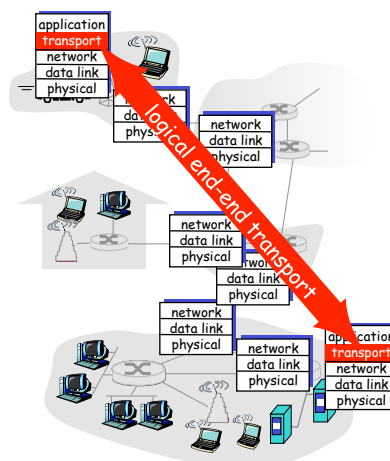
Household analogy:

3 kids sending letters to 3 other kids

- ❑ processes = kids
- ❑ app messages = letters in envelopes
- ❑ hosts = houses
- ❑ transport protocol = parents
- ❑ network-layer protocol = postal service

Internet transport-layer protocols

- ❑ reliable, in-order delivery (TCP)
 - ❖ congestion control
 - ❖ flow control
 - ❖ connection setup
- ❑ unreliable, unordered delivery: UDP
 - ❖ no-frills extension of "best-effort" IP
- ❑ services not available:
 - ❖ delay guarantees
 - ❖ bandwidth guarantees



Outline

- ❑ 3.1 Transport-layer services
- ❑ 3.2 Multiplexing and demultiplexing
- ❑ 3.3 Connectionless transport: UDP
- ❑ 3.4 Principles of reliable data transfer
- ❑ 3.5 Connection-oriented transport: TCP
- ❑ 3.6 TCP congestion control

Multiplexing/demultiplexing


Demultiplexing at rcv host:

delivering received segments to correct socket

Multiplexing at send host:

gathering data from multiple sockets, enveloping data with header (later used for demultiplexing)

■ = socket ○ = process




How demultiplexing works

□ host receives IP datagrams

- ❖ each datagram has source IP address, destination IP address
- ❖ each datagram carries 1 transport-layer segment
- ❖ each segment has source, destination port number

□ host uses IP addresses & port numbers to direct segment to appropriate socket



TCP/UDP segment format

Connectionless demultiplexing

□ Create sockets with port numbers:

```
DatagramSocket mySocket1 = new  
    DatagramSocket(12534);
```

```
DatagramSocket mySocket2 = new  
    DatagramSocket(12535);
```

□ UDP socket identified by two-tuple:

(dest IP address, dest port number)


□ When host receives UDP segment:

- ❖ checks destination port number in segment
- ❖ directs UDP segment to socket with that port number

□ IP datagrams with different source IP addresses and/or source port numbers can be directed to same socket

Connectionless demux (cont)

```
DatagramSocket serverSocket = new DatagramSocket(6428);
```



SP provides "return address"

University College Cork CS2505: Transport Layer

3-11

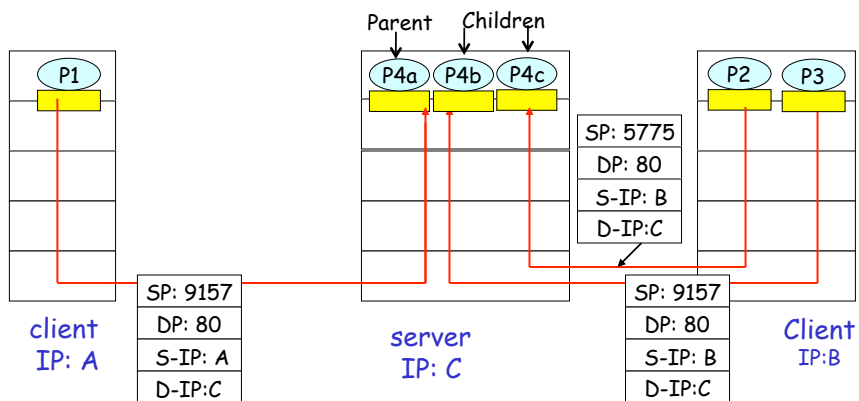
Connection-oriented demux

- ❑ TCP socket identified by 4-tuple:
 - ❖ source IP address
 - ❖ source port number
 - ❖ dest IP address
 - ❖ dest port number
- ❑ receiving host uses all four values to direct segment to appropriate socket
- ❑ Server host may support many simultaneous TCP sockets:
 - ❖ each socket identified by its own 4-tuple
- ❑ Web servers have different sockets for each connecting client
 - ❖ non-persistent HTTP will have different socket for each request

University College Cork CS2505: Transport Layer

3-12

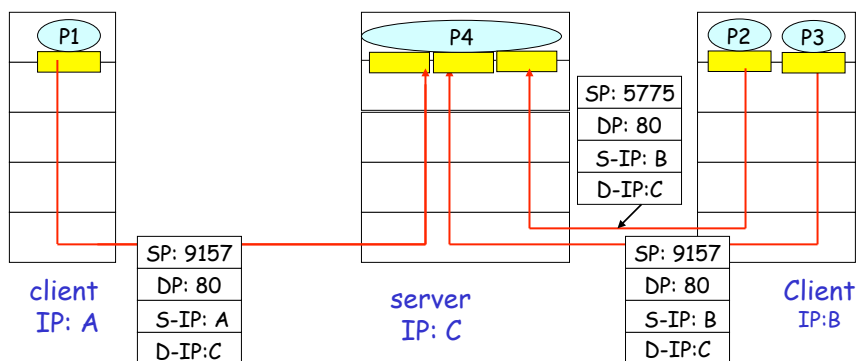
Connection-oriented demux (cont)



University College Cork CS2505: Transport Layer

3-13

Connection-oriented demux: Threaded Web Server



University College Cork CS2505: Transport Layer

3-14

Outline

- ❑ 3.1 Transport-layer services
- ❑ 3.2 Multiplexing and demultiplexing
- ❑ 3.3 Connectionless transport: UDP
- ❑ 3.4 Principles of reliable data transfer
- ❑ 3.5 Connection-oriented transport: TCP
- ❑ 3.6 TCP congestion control

UDP: User Datagram Protocol [RFC 768]

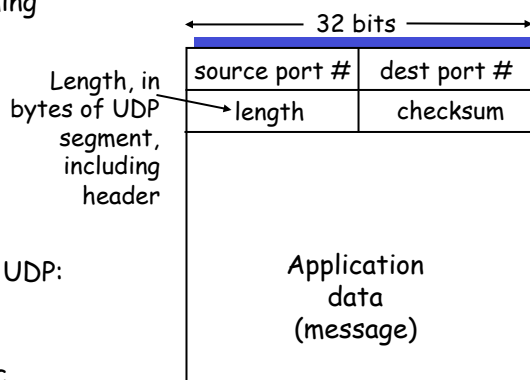
- ❑ "no frills," "bare bones" Internet transport protocol
- ❑ "best effort" service, UDP segments may be:
 - ❖ lost
 - ❖ delivered out of order to app
- ❑ **connectionless:**
 - ❖ no handshaking between UDP sender, receiver
 - ❖ each UDP segment handled independently of others

Why is there a UDP?

- ❑ no connection establishment (which can add delay)
- ❑ no (delay for) recovering lost segments as in TCP
- ❑ simple: no connection state at sender, receiver
- ❑ small segment header
- ❑ no congestion control: UDP can blast away as fast as desired

UDP: more

- ❑ often used for streaming multimedia apps
 - ❖ loss tolerant
 - ❖ rate sensitive
- ❑ other UDP uses
 - ❖ DNS
 - ❖ SNMP
- ❑ reliable transfer over UDP: add reliability at application layer
 - ❖ application-specific error recovery!



UDP segment format

UDP checksum

Goal: detect "errors" (e.g., flipped bits) in transmitted segment

Sender:

- ❑ treat segment contents as sequence of 16-bit integers
- ❑ checksum: addition (1's complement sum) of segment contents
- ❑ sender puts checksum value into UDP checksum field

Receiver:

- ❑ compute checksum of received segment
- ❑ check if computed checksum equals checksum field value:
 - ❖ NO - error detected
 - ❖ YES - no error detected. *But maybe errors nonetheless?*

Internet Checksum Example

□ Note

- ❖ When adding numbers, a carryout from the most significant bit needs to be added to the result

□ Example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

University College Cork CS2505: Transport Layer

3-19

Outline


- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
- 3.6 TCP congestion control

University College Cork CS2505: Transport Layer

3-20

Principles of Reliable data transfer

- important in app., transport, link layers
- top-10 list of important networking topics!




(a) provided service

- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Principles of Reliable data transfer

- important in app., transport, link layers
- top-10 list of important networking topics!



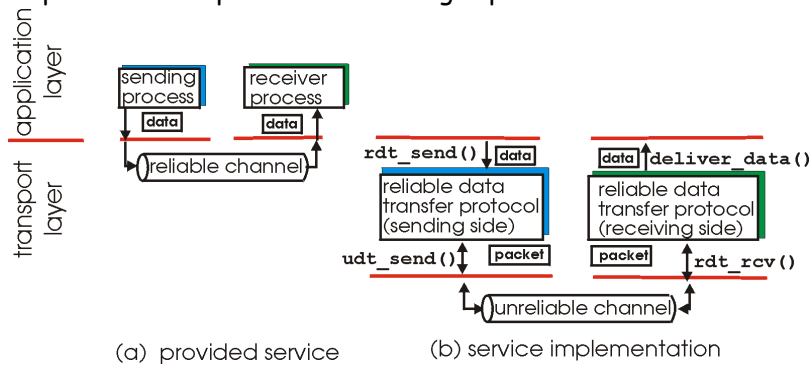
(a) provided service

(b) service implementation

- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

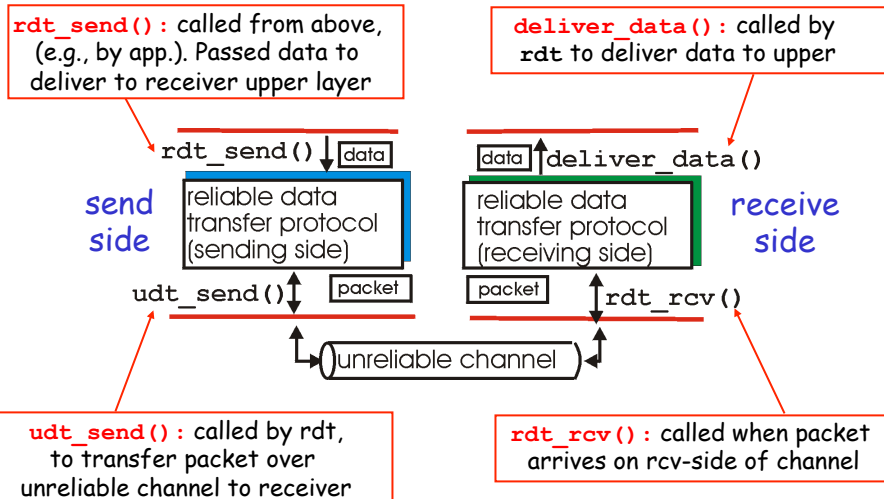
Principles of Reliable data transfer

- important in app., transport, link layers
- top-10 list of important networking topics!



- characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Reliable data transfer: getting started

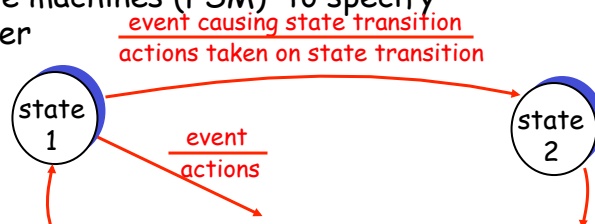


Reliable data transfer: getting started

In this section we will:

- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- For simplicity assume consider "data" in one direction only and (initially) no out-of-order delivery
- use finite state machines (FSM) to specify sender, receiver

state: when in this "state" next state uniquely determined by next event




University College Cork CS2505: Transport Layer

3-25

Rdt1.0: reliable transfer over a reliable channel

- underlying channel perfectly reliable
 - ❖ no bit errors
 - ❖ no loss of packets
- separate FSMs for sender, receiver:
 - ❖ sender sends data into underlying channel
 - ❖ receiver read data from underlying channel



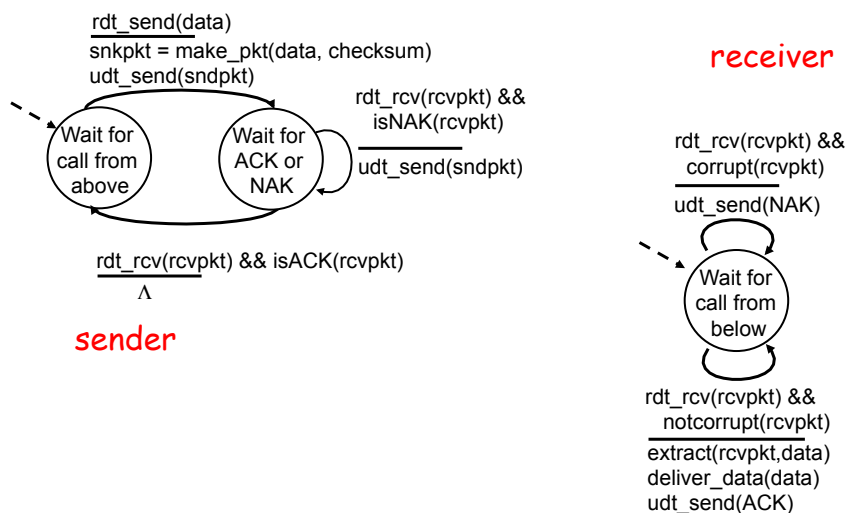
University College Cork CS2505: Transport Layer

3-26

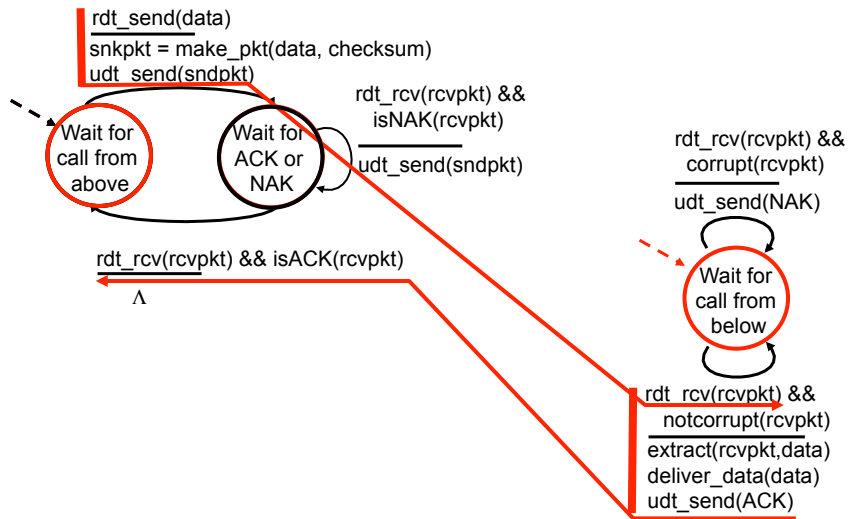
Rdt2.0: channel with bit errors

- ❑ underlying channel may flip bits in packet
 - ❖ checksum to detect bit errors
- ❑ *the question*: how to recover from errors:
 - ❖ *acknowledgements (ACKs)*: receiver explicitly tells sender that pkt received OK
 - ❖ *negative acknowledgements (NAKs)*: receiver explicitly tells sender that pkt had errors
 - ❖ sender retransmits pkt on receipt of NAK
- ❑ new mechanisms in rdt2.0 (beyond rdt1.0):
 - ❖ error detection
 - ❖ receiver feedback: control msgs (ACK,NAK) rcvr→sender

rdt2.0: FSM specification



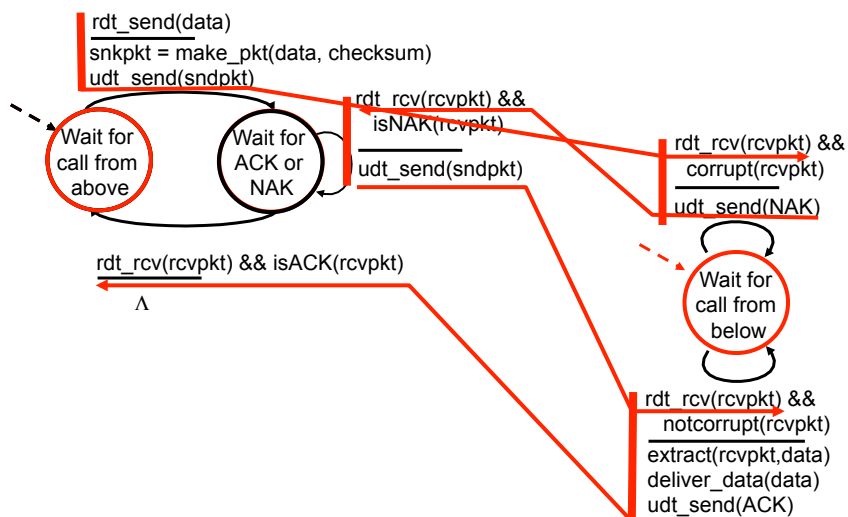
rdt2.0: operation with no errors



University College Cork CS2505: Transport Layer

3-29

rdt2.0: error scenario



University College Cork CS2505: Transport Layer

3-30

rdt2.0 has a fatal flaw!

What happens if ACK/NAK corrupted?

- sender doesn't know what happened at receiver!
- can't just retransmit: possible duplicate

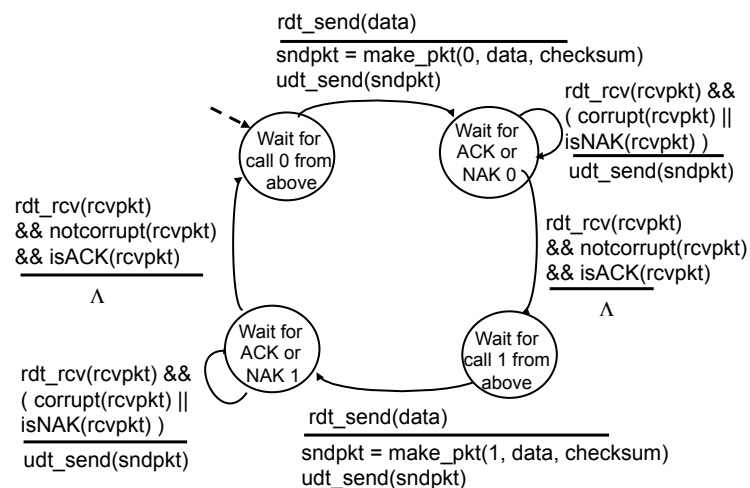
Handling duplicates:

- sender retransmits current pkt if ACK/NAK garbled
- sender adds *sequence number* to each pkt
- receiver discards (doesn't deliver up) duplicate pkt

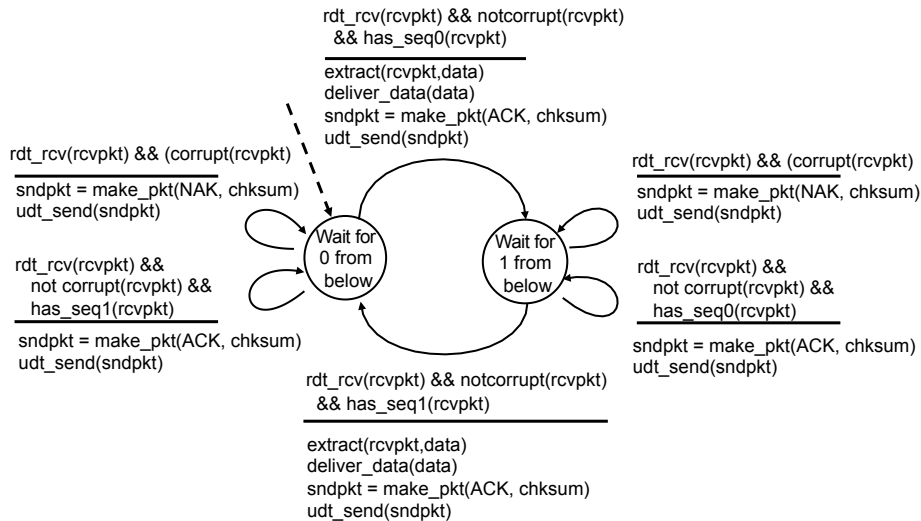
stop and wait

Sender sends one packet, then waits for receiver Response before sending anything

rdt2.1: sender, handles garbled ACK/NAKs



rdt2.1: receiver, handles garbled ACK/NAKs



University College Cork CS2505: Transport Layer

3-33

rdt2.1: discussion

Sender:

- ❑ seq # added to pkt
- ❑ two seq. #'s (0,1) will suffice. Why?
- ❑ must check if received ACK/NAK corrupted
- ❑ twice as many states
 - ❖ state must "remember" whether "current" pkt has 0 or 1 seq. #

Receiver:

- ❑ must check if received packet is duplicate
 - ❖ state indicates whether 0 or 1 is expected pkt seq #
- ❑ note: receiver can *not* know if its last ACK/NAK received OK at sender

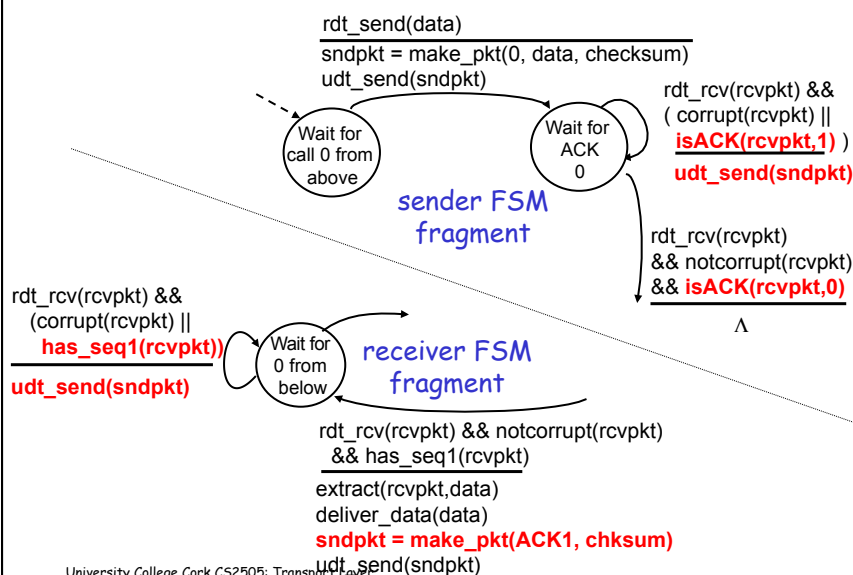
University College Cork CS2505: Transport Layer

3-34

rdt2.2: a NAK-free protocol

- ❑ same functionality as rdt2.1, using ACKs only
- ❑ instead of NAK, receiver sends ACK for last pkt received OK
 - ❖ receiver must *explicitly* include seq # of pkt being ACKed
 - ❖ sender then knows that the current packet was not received correctly
- ❑ duplicate ACK at sender results in same action as NAK: *retransmit current pkt*
- ❑ This is a simpler protocol because it does away with NAKs

rdt2.2: sender, receiver fragments



rdt3.0: channels with errors and loss

New assumption:

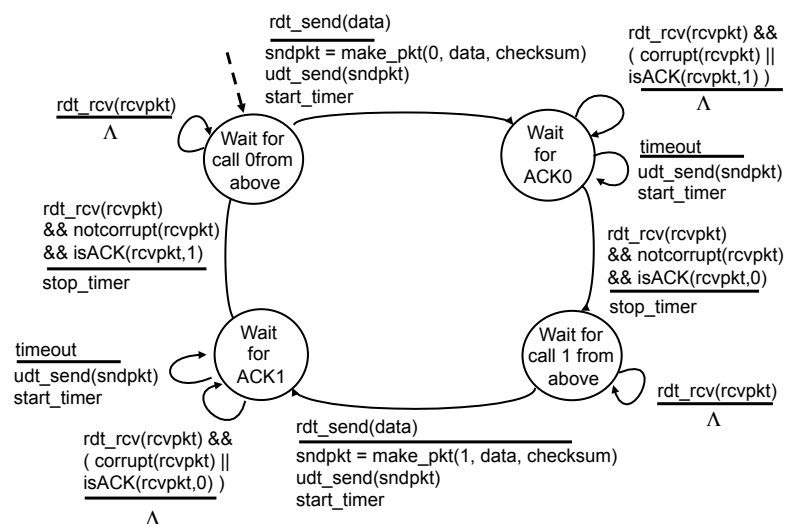
underlying channel can also lose packets (data or ACKs)

- ❖ checksum, seq. #, ACKs, retransmissions will be of help, but not enough



Approach: sender waits "reasonable" amount of time for ACK

- ❑ retransmits if no ACK received in this time
- ❑ if pkt (or ACK) just delayed (not lost):
 - ❖ retransmission will be duplicate, but use of seq. #'s already handles this
 - ❖ receiver must specify seq # of pkt being ACKed
- ❑ requires countdown timer

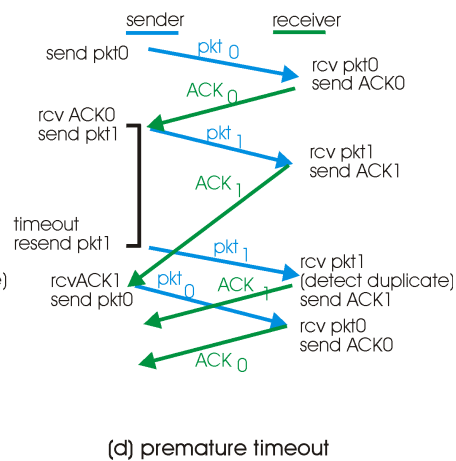
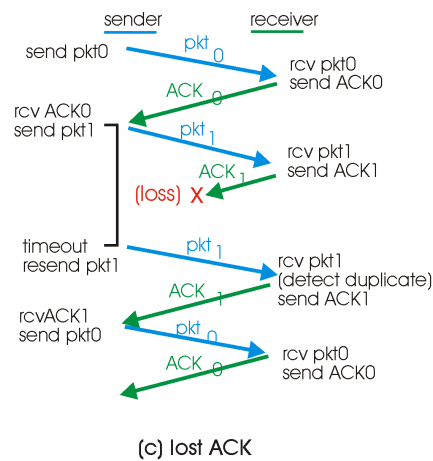
rdt3.0 sender



rdt3.0 in action



rdt3.0 in action



Performance of rdt3.0

- rdt3.0 works, but performance stinks
- eg: 1 Gb/s link, 15 ms propagation delay, 8000 bit packet:


$$d_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ b/s}} = 8 \text{ microseconds}$$

- ❖ U_{sender} : **utilization** - fraction of time sender busy sending

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- ❖ 1KB pkt every 30 msec → 33KB/sec throughput over 1 Gb/s link
- ❖ network protocol limits use of physical resources!

rdt3.0: stop-and-wait operation

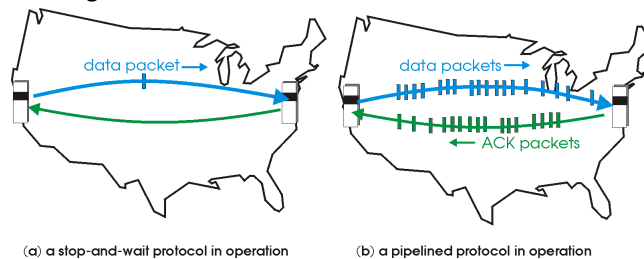


$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

Pipelined protocols

Pipelining: sender allows multiple, "in-flight", yet-to-be-acknowledged pkts

- ❖ range of sequence numbers must be increased
- ❖ buffering at sender and/or receiver




(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

- ❑ Two generic forms of pipelined protocols: *go-Back-N*, *selective repeat*

Pipelining: increased utilization



Pipelining Protocols

Go-back-N: overview

- **sender:** up to N unACKed pkts in pipeline
- **receiver:** only sends cumulative ACKs
 - ❖ doesn't ACK pkt if there's a gap
- **sender:** has timer for oldest unACKed pkt
 - ❖ if timer expires: retransmit all unACKed packets


Selective Repeat: overview

- **sender:** up to N unACKed packets in pipeline
- **receiver:** ACKs individual pkts
- **sender:** maintains timer for each unACKed pkt
 - ❖ if timer expires: retransmit only unACKed packet

Go-Back-N

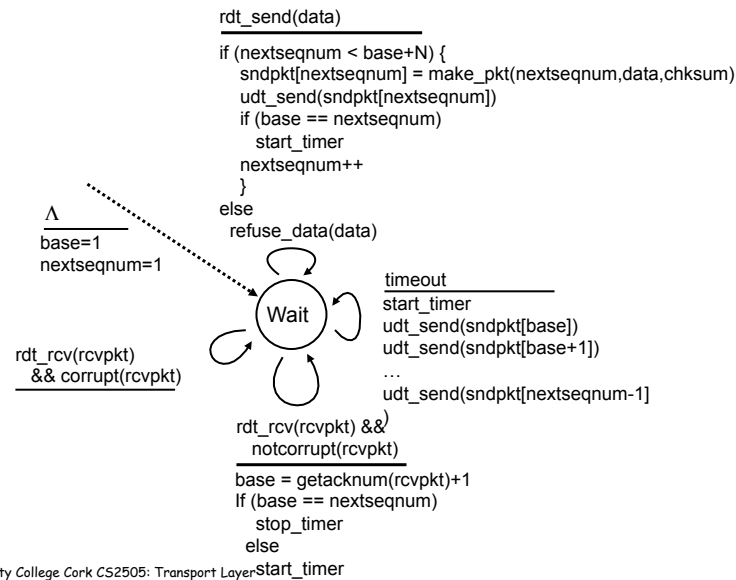
Sender:

- k-bit seq # in pkt header
- "sliding window" of up to N, consecutive unACKed pkts allowed




- **ACK(n):** ACKs all pkts up to, including seq # n - "cumulative ACK"
 - ❖ may receive duplicate ACKs (see receiver)
- **timeout(n):** retransmit pkt n and all higher seq # pkts in window

GBN: sender extended FSM



3-47

GBN: receiver extended FSM

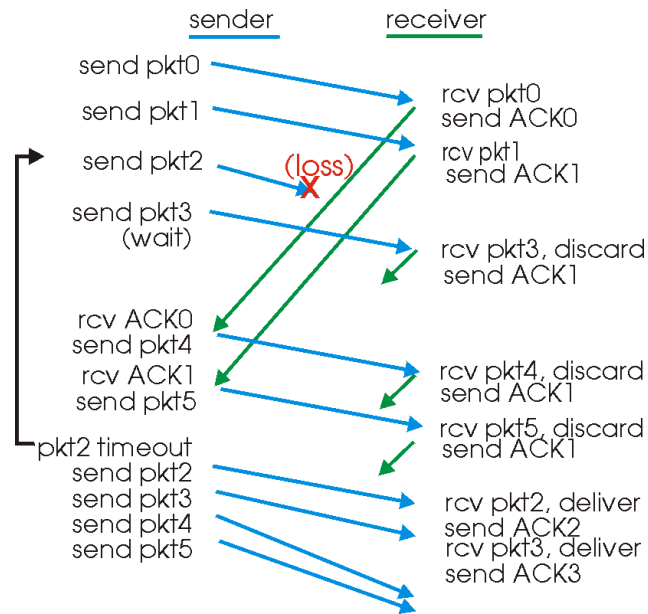


- ❑ **ACK-only:** always send ACK for correctly-received pkt with highest *in-order* seq #
 - ❖ may generate duplicate ACKs
 - ❖ need only remember `expectedseqnum`
- ❑ **out-of-order pkt:**
 - ❖ discard (don't buffer) -> **no receiver buffering!**
 - ❖ Re-ACK pkt with highest in-order seq #

University College Cork CS2505: Transport Layer

3-48

GBN in action



University College Cork CS2505: Transport Layer

3-49

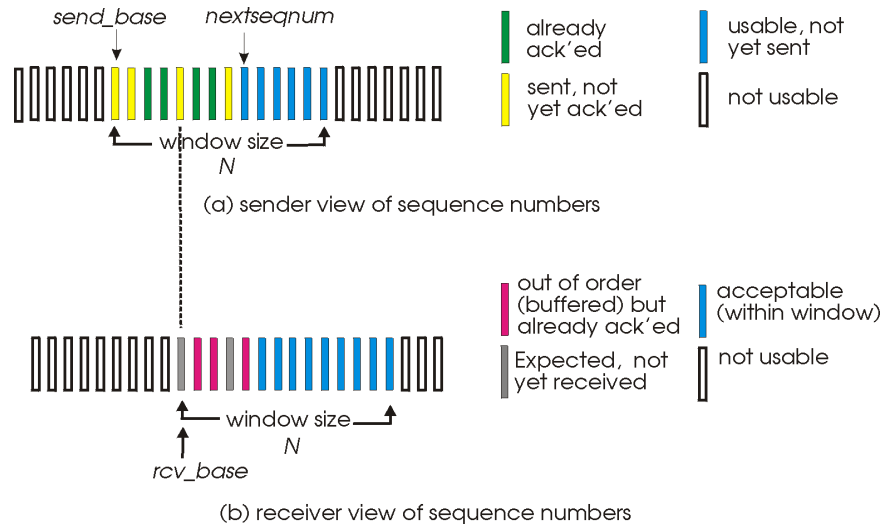
Selective Repeat

- ❑ Go-back-N can be inefficient if there can be many pkts in pipeline and an error occurs
 - ❖ All these packets will be retransmitted unnecessarily
- ❑ With selective repeat receiver *individually* acknowledges all correctly received pkts
 - ❖ buffers pkts, as needed, for eventual in-order delivery to upper layer
 - ❖ sender only resends pkts for which ACK not received
 - sender timer for each unACKed pkt
 - ❖ sender window
 - N consecutive seq #s
 - again limits seq #s of sent, unACKed pkts

University College Cork CS2505: Transport Layer

3-50

Selective repeat: sender, receiver windows



University College Cork CS2505: Transport Layer

3-51

Selective repeat

sender

data from above :

- if next available seq # in window, send pkt

timeout(n):

- resend pkt n , restart timer

ACK(n) in $[sendbase, sendbase+N]$:

- mark pkt n as received
- if n smallest unACKed pkt, advance window base to next unACKed seq #

receiver

pkt n in $[rcvbase, rcvbase+N-1]$

- send ACK(n)
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt

pkt n in $[rcvbase-N, rcvbase-1]$

- ACK(n)

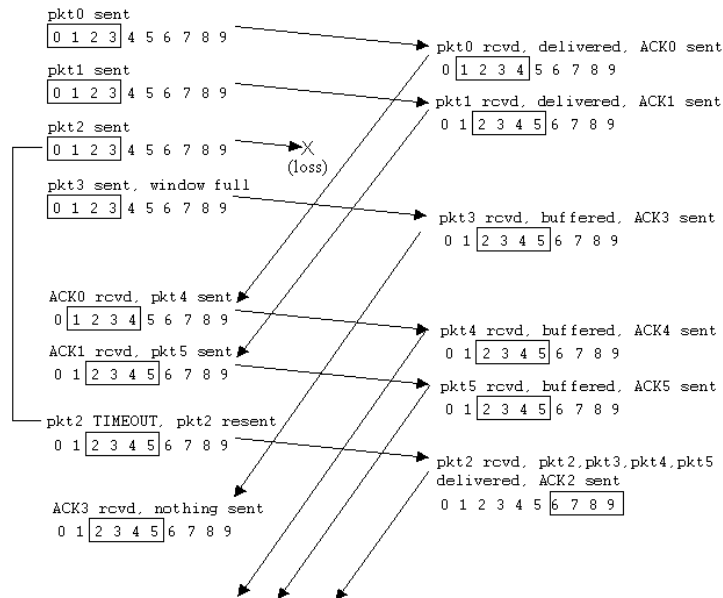
otherwise:

- ignore

University College Cork CS2505: Transport Layer

3-52

Selective repeat in action



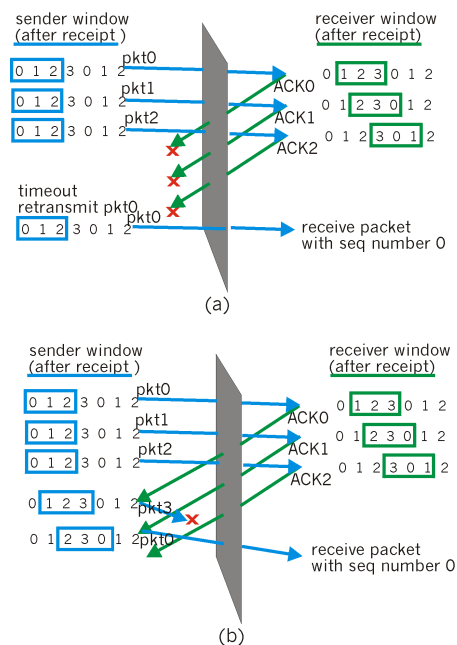
3-53

Selective repeat: dilemma

Example:

- ❑ seq #'s: 0, 1, 2, 3
- ❑ window size=3
- ❑ receiver sees no difference in two scenarios!
- ❑ incorrectly passes duplicate data as new in (a)

Q: what relationship between seq # size and window size?



3-54

Outline

- ❑ 3.1 Transport-layer services
- ❑ 3.2 Multiplexing and demultiplexing
- ❑ 3.3 Connectionless transport: UDP
- ❑ 3.4 Principles of reliable data transfer
- ❑ 3.5 Connection-oriented transport: TCP
 - ❖ basics
 - ❖ reliable data transfer
 - ❖ flow control
- ❑ 3.6 TCP congestion control

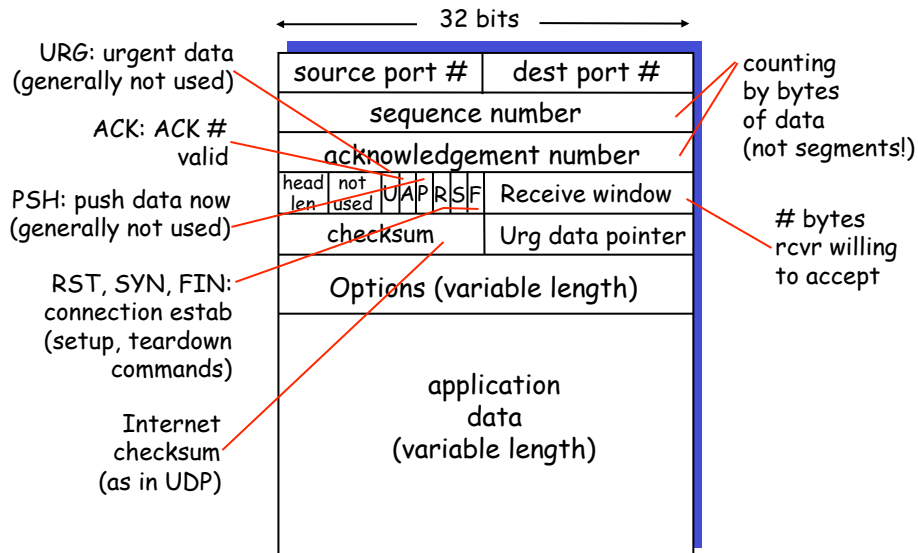
TCP: Overview

RFCs: 793, 1122, 1323, 2018, 2581

- ❑ **point-to-point:**
 - ❖ one sender, one receiver
- ❑ **reliable, in-order byte stream:**
 - ❖ no "message boundaries"
- ❑ **pipelined:**
 - ❖ TCP congestion and flow control set window size
- ❑ **send & receive buffers**
- ❑ **full duplex data:**
 - ❖ bi-directional data flow in same connection
 - ❖ MSS: maximum segment size
- ❑ **connection-oriented:**
 - ❖ handshaking (exchange of control msgs) init's sender, receiver state before data exchange
- ❑ **flow controlled:**
 - ❖ sender will not overwhelm receiver



TCP segment structure



University College Cork CS2505: Transport Layer

3-57

TCP Connection Management

Recall: TCP sender, receiver establish "connection" before exchanging data segments

- initialize TCP variables:
 - ❖ seq. #s
 - ❖ buffers, flow control info (e.g. `RcvWindow`)
- *client*: connection initiator


```
Socket clientSocket = new
Socket("hostname", "port
number");
```
- *server*: contacted by client


```
Socket connectionSocket =
welcomeSocket.accept();
```

Three way handshake:

Step 1: client host sends TCP SYN segment to server

- ❖ specifies initial seq #
- ❖ no data

Step 2: server host receives SYN, replies with SYNACK segment

- ❖ server allocates buffers
- ❖ specifies server initial seq. #

Step 3: client receives SYNACK, replies with ACK segment, which may contain data

University College Cork CS2505: Transport Layer

3-58


TCP Connection Management (cont.)

Closing a connection:

client closes socket:
`clientSocket.close();`

Step 1: client end system
sends TCP FIN control
segment to server

Step 2: server receives
FIN, replies with ACK.
Closes connection, sends
FIN.




TCP Connection Management (cont.)

Step 3: client receives FIN,
replies with ACK.

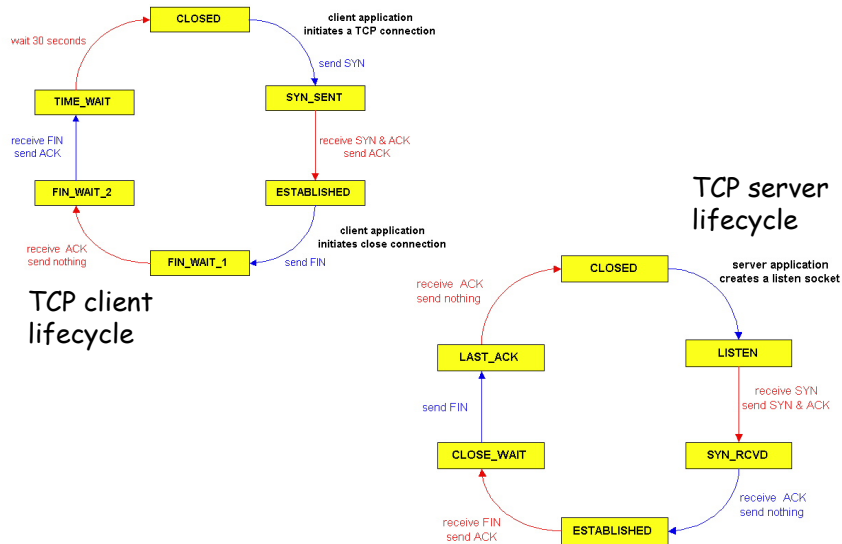
- ❖ Enters "timed wait" -
will respond with ACK
to received FINs

Step 4: server, receives
ACK. Connection closed.

Note: with small
modification, can handle
simultaneous FINs.



TCP Connection Management (cont)



University College Cork CS2505: Transport Layer

3-61


TCP seq. #'s and ACKs

Seq. #'s:

- ❖ byte stream "number" of first byte in segment's data

ACKs:

- ❖ seq # of next byte expected from other side
- ❖ cumulative ACK




University College Cork CS2505: Transport Layer

3-62

Example TCP Session

Starting sequence numbers:

- ❖ increase by one during connection opening
- ❖ chosen randomly (shown as zero in figure)
- ❖ No increase in seqno for empty ACK



University College Cork CS2505: Transport Layer

3-63

TCP reliable data transfer

- ❑ TCP creates rdt service on top of IP's unreliable service
- ❑ pipelined segments
- ❑ cumulative ACKs
- ❑ TCP uses single retransmission timer
- ❑ retransmissions are triggered by:
 - ❖ timeout events
 - ❖ duplicate ACKs
- ❑ initially consider simplified TCP sender:
 - ❖ ignore duplicate ACKs
 - ❖ ignore flow control, congestion control

University College Cork CS2505: Transport Layer

3-64

TCP sender events:

data rcvd from app:

- ❑ create segment with seq #
- ❑ seq # is byte-stream number of first data byte in segment
- ❑ start timer if not already running (think of timer as for oldest unACKed segment)
- ❑ expiration interval:
TimeOutInterval

timeout:

- ❑ retransmit segment that caused timeout
- ❑ restart timer

ACK rcvd:

- ❑ if acknowledges previously unACKed segments
 - ❖ update what is known to be ACKed
 - ❖ start timer if there are outstanding segments

```
NextSeqNum = InitialSeqNum
SendBase = InitialSeqNum

loop (forever) {
  switch(event)

    event: data received from application above
    create TCP segment with sequence number NextSeqNum
    if (timer currently not running)
      start timer
    pass segment to IP
    NextSeqNum = NextSeqNum + length(data)

    event: timer timeout
    retransmit not-yet-acknowledged segment with
      smallest sequence number
    start timer

    event: ACK received, with ACK field value of y
    if (y > SendBase) {
      SendBase = y
      if (there are currently not-yet-acknowledged segments)
        start timer
    }

} /* end of loop forever */
```

TCP sender (simplified)

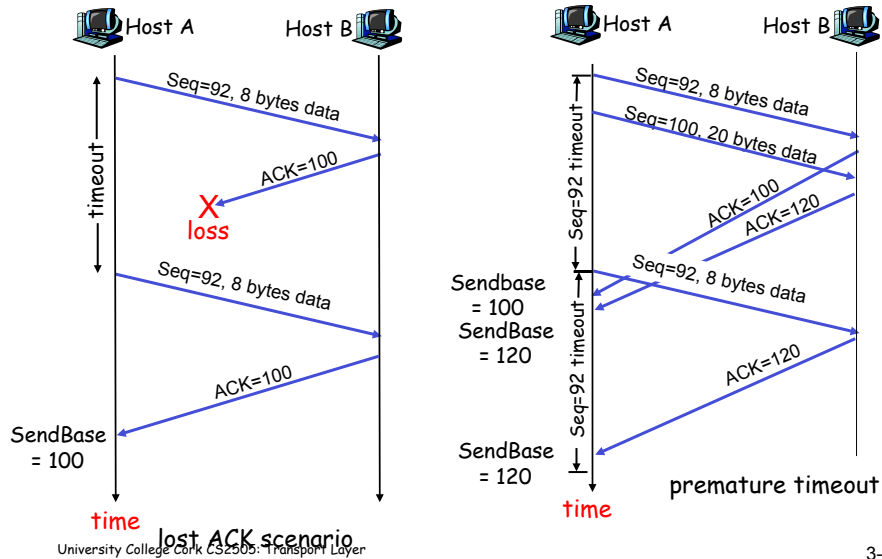
Comment:

- SendBase-1: last cumulatively ACKed byte

Example:


- SendBase-1 = 71; y = 73, so the rcvr wants 73+ ; y > SendBase, so that new data is ACKed

TCP: retransmission scenarios



3-67

TCP retransmission scenarios (more)



University College Cork CS2505: Transport Layer

3-68

TCP ACK generation [RFC 1122, RFC 2581]

Event at Receiver	TCP Receiver action
Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
Arrival of in-order segment with expected seq #. One other segment has ACK pending	Immediately send single cumulative ACK, ACKing both in-order segments
Arrival of out-of-order segment higher-than-expected seq. # . Gap detected	Immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
Arrival of segment that partially or completely fills gap	Immediate send ACK, provided that segment starts at lower end of gap

University College Cork CS2505: Transport Layer

3-69

TCP Selective ACKs [RFC 2018]


- ❑ A non-mandatory extension to TCP cumulative ACKs that is widely used
- ❑ Selective ACK (SACK) allows receiver to ACK a sequence of bytes in addition to number of next expected byte
- ❑ Use of SACK is negotiated during TCP connection opening
 - ❖ uses TCP options field to convey sequence number ranges

University College Cork CS2505: Transport Layer

3-70

Fast Retransmit

- time-out period often relatively long:
 - ❖ long delay before resending lost packet
- detect lost segments via duplicate ACKs.
 - ❖ sender often sends many segments back-to-back
 - ❖ if segment is lost, there will likely be many duplicate ACKs for that segment
- If sender receives 3 ACKs for same data, it assumes that segment after ACKed data was lost:
 - ❖ fast retransmit: resend segment before timer expires



Fast retransmit algorithm:

```
event: ACK received, with ACK field value of y
    if (y > SendBase) {
        SendBase = y
        if (there are currently not-yet-acknowledged segments)
            start timer
    }
    else {
        increment count of dup ACKs received for y
        if (count of dup ACKs received for y = 3) {
            resend segment with sequence number y
        }
    }
```

a duplicate ACK for
already ACKed segment

fast retransmit

TCP Round Trip Time and Timeout

Q: how to set TCP
timeout value?

- ❑ longer than RTT
 - ❖ but RTT varies
- ❑ too short: premature timeout
 - ❖ unnecessary retransmissions
- ❑ too long: slow reaction to segment loss

Q: how to estimate RTT?

- ❑ **SampleRTT**: measured time from segment transmission until ACK receipt
 - ❖ ignore retransmissions
- ❑ **SampleRTT** will vary, want estimated RTT "smoother"
 - ❖ average several recent measurements, not just current **SampleRTT**


TCP Round Trip Time and Timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- Exponential weighted moving average
- influence of past sample decreases exponentially fast
- typical value: $\alpha = 0.125$

Example RTT estimation:

RTT: gaia.cs.umass.edu to fantasia.eurecom.fr



TCP Round Trip Time and Timeout

Setting the timeout

- EstimatedRTT plus "safety margin"
 - ❖ large variation in EstimatedRTT -> larger safety margin
- first estimate of how much SampleRTT deviates from EstimatedRTT:

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$


(typically, $\beta = 0.25$)

Then set timeout interval:

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$

TCP Flow Control

- receive side of TCP connection has a receive buffer:




flow control

sender won't overflow receiver's buffer by transmitting too much, too fast

- *speed-matching service*: matching send rate to receiving application's drain rate
- app process may be slow at reading from buffer

TCP Flow control: how it works



(suppose TCP receiver discards out-of-order segments)

□ unused buffer space:

= $rwnd$

= $RcvBuffer - [LastByteRcvd - LastByteRead]$

□ receiver: advertises unused buffer space by including $rwnd$ value in segment header

□ sender: limits # of unACKed bytes to $rwnd$
 ❖ guarantees receiver's buffer doesn't overflow

TCP Flow Control Example

□ Example: slow receiver

- ❖ Recv buffer fills up and window shrinks to 0
- ❖ Send TCP learns of empty window and stops
- ❖ Send buffer fills up with bytes from appl process
- ❖ Send TCP asks OS to block sender appl process

□ Once receiver catches up

- ❖ Window opens, Send TCP learns new window size
- ❖ Send TCP resumes transmission
- ❖ Send TCP buffer frees up
- ❖ Send TCP asks OS to unblock sender process

Outline


- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
- 3.6 TCP congestion control

TCP congestion control:

- **goal:** TCP sender should transmit as fast as possible, but without congesting network
 - ❖ **Q:** how to find rate *just* below congestion level
- decentralized: each TCP sender sets its own rate, based on **implicit** feedback:
 - ❖ **ACK:** segment received (a good thing!), network not congested, so increase sending rate
 - ❖ **lost segment:** assume loss due to congested network, so decrease sending rate

TCP congestion control: bandwidth probing

- "probing for bandwidth": increase transmission rate on receipt of ACK, until eventually loss occurs, then decrease transmission rate
 - ❖ continue to increase on ACK, decrease on loss (since available bandwidth is changing, depending on other connections in network)



- Q: how fast to increase/decrease?

- ❖ details to follow

University College Cork CS2505: Transport Layer

3-83

TCP Congestion Control: details

- sender limits rate by limiting number of unACKed bytes "in pipeline":


$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- ❖ cwnd: differs from rwnd (how, why?)
- ❖ sender limited by $\min(\text{cwnd}, \text{rwnd})$

- roughly,

$$\text{rate} = \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

- cwnd is dynamic, function of perceived network congestion



University College Cork CS2505: Transport Layer

3-84

TCP Congestion Control: more details

segment loss event: reducing cwnd


- ❑ timeout: no response from receiver
 - ❖ cut **cwnd** to 1
- ❑ 3 duplicate ACKs: at least some segments getting through (recall fast retransmit)
 - ❖ cut **cwnd** in half, less aggressively than on timeout

ACK received: increase cwnd

- ❑ slowstart phase:
 - ❖ increase exponentially fast (despite name) at connection start, or following timeout
- ❑ congestion avoidance:
 - ❖ increase linearly

TCP Slow Start

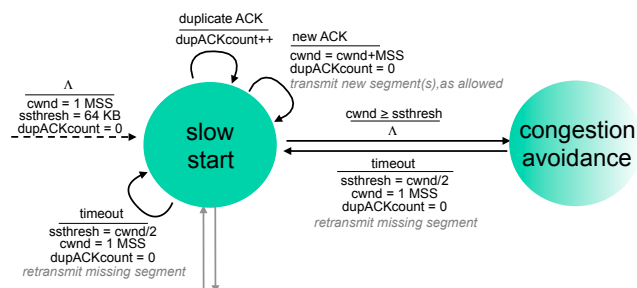
- ❑ when connection begins, **cwnd** = 1 MSS
 - ❖ example: MSS = 500 bytes & RTT = 200 msec
 - ❖ initial rate = 20 kbps
- ❑ available bandwidth may be \gg MSS/RTT
 - ❖ desirable to quickly ramp up to respectable rate
- ❑ increase rate exponentially until first loss event or when threshold reached
 - ❖ double **cwnd** every RTT
 - ❖ done by incrementing **cwnd** by 1 for every ACK received



Transitioning into/out of slowstart

ssthresh: cwnd threshold maintained by TCP

- ❑ on loss event: set **ssthresh** to **cwnd/2**
 - ❖ remember (half of) TCP rate when congestion last occurred
- ❑ when **cwnd** \geq **ssthresh**: transition from slowstart to congestion avoidance phase



University College Cork CS2505: Transport Layer

3-87

TCP: congestion avoidance

- ❑ when **cwnd** $>$ **ssthresh**
 - grow **cwnd** linearly
 - ❖ increase **cwnd** by 1 MSS per RTT
 - ❖ approach possible congestion slower than in slowstart
 - ❖ implementation: $cwnd = cwnd + MSS/cwnd$ for each ACK received

AIMD

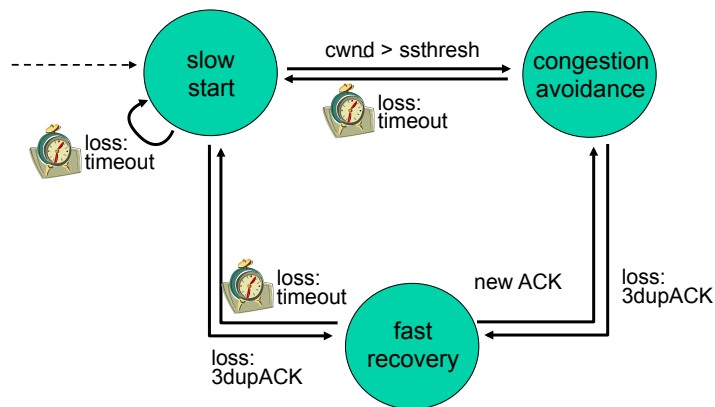
- ❖ **ACKs**: increase **cwnd** by 1 MSS per RTT: additive increase
- ❖ **loss**: cut **cwnd** in half (non-timeout-detected loss): multiplicative decrease

AIMD: Additive Increase
Multiplicative Decrease

University College Cork CS2505: Transport Layer

3-88

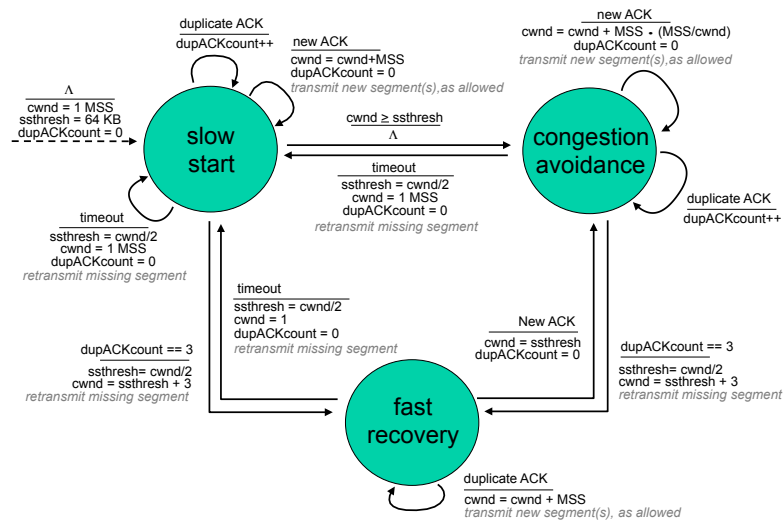
TCP congestion control FSM: overview



University College Cork CS2505: Transport Layer

3-89


TCP congestion control FSM: details



University College Cork CS2505: Transport Layer

3-90

Popular "flavours" of TCP



University College Cork CS2505: Transport Layer

3-91

Summary: TCP Congestion Control

- when $cwnd < sssthresh$, sender in **slow-start** phase, window grows exponentially.
- when $cwnd \geq sssthresh$, sender is in **congestion-avoidance** phase, window grows linearly.
- when **triple duplicate ACK** occurs, $sssthresh$ set to $cwnd/2$, $cwnd$ set to $\sim sssthresh$
- when **timeout** occurs, $sssthresh$ set to $cwnd/2$, $cwnd$ set to 1 MSS.

University College Cork CS2505: Transport Layer

3-92

Summary

- ❑ principles behind transport layer services:
 - ❖ multiplexing, demultiplexing
 - ❖ reliable data transfer
 - ❖ flow control
 - ❖ congestion control
- ❑ instantiation and implementation in the Internet
 - ❖ UDP
 - ❖ TCP